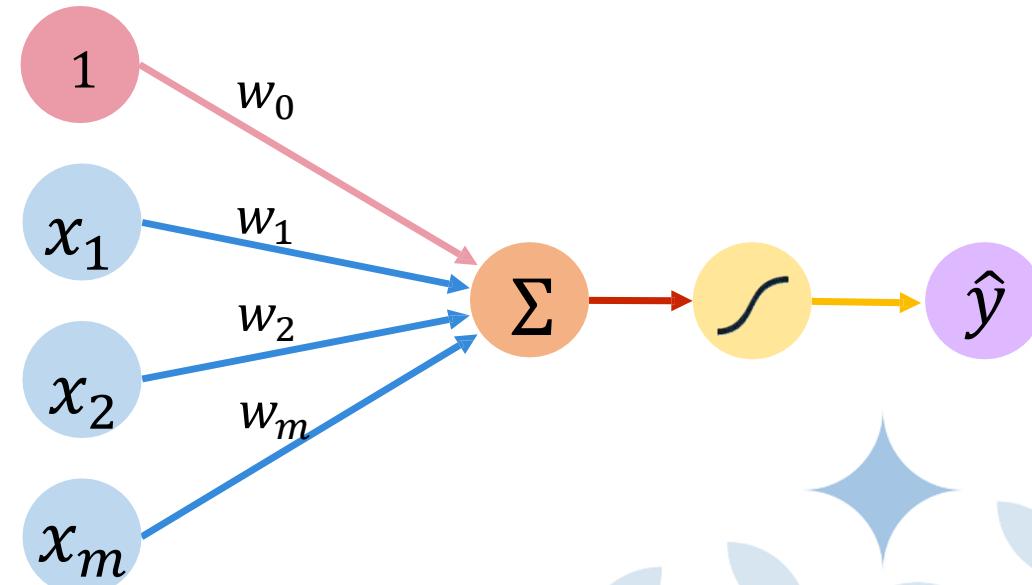
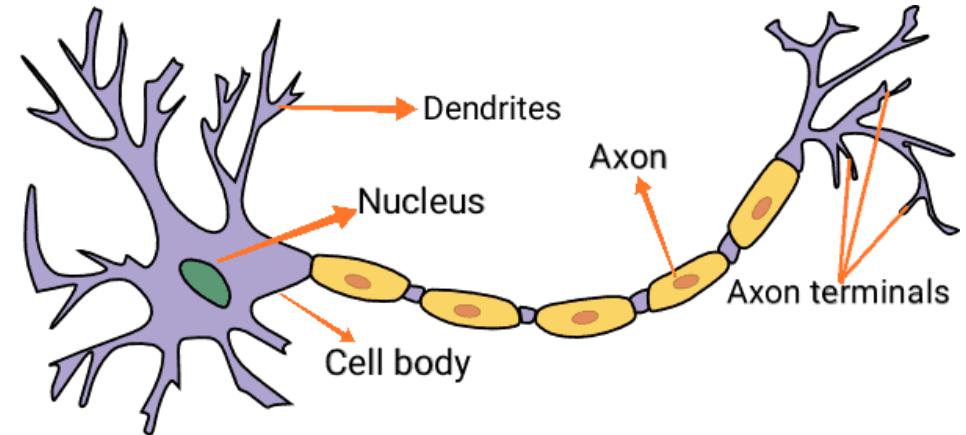


# Artificial Neural Networks

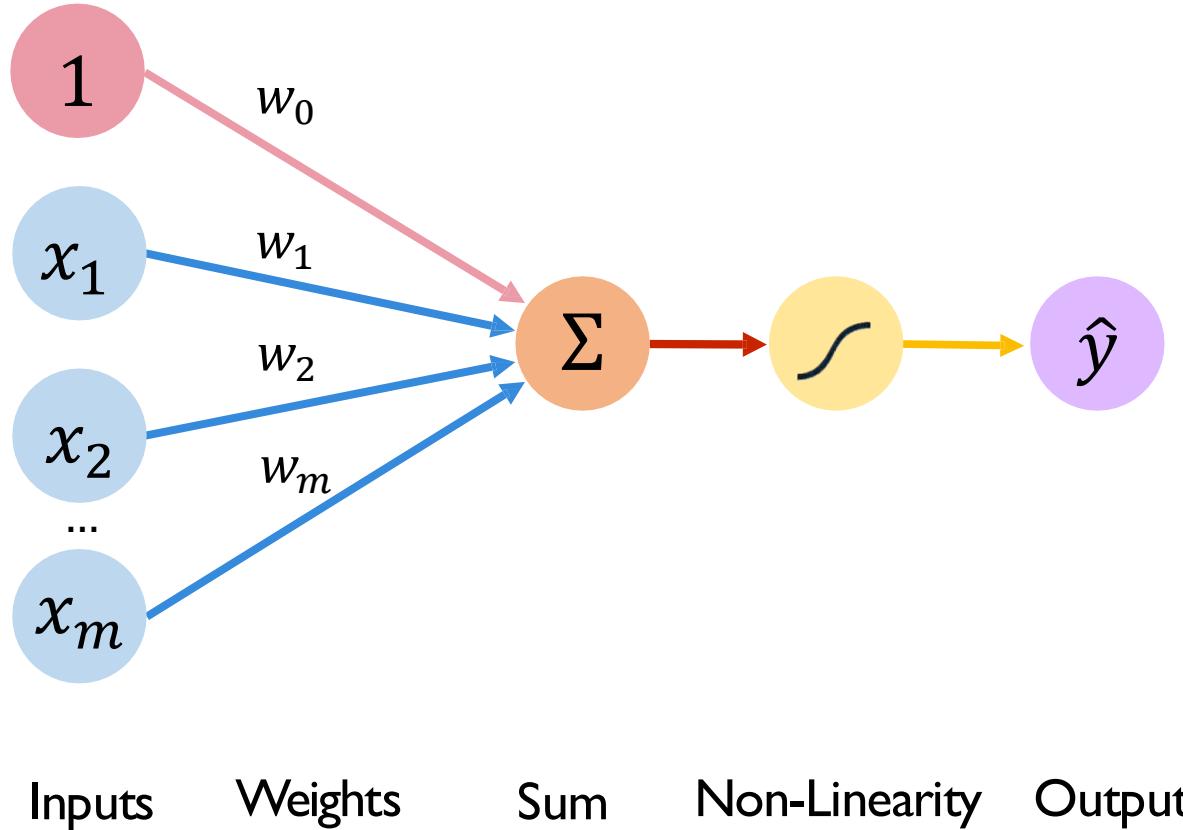
**Qiang Sun**  
**University of Toronto**

# Artificial Neural Network

- A network of **artificial** neurons that **mimics real biological** neural networks
  - Each has a body, axon, and many dendrites
  - A neuron can fire or rest
  - A threshold to activate (fire)
- **Perceptron** (Frank Rosenblatt, 1956)
  - Fundamental building block of deep learning
  - What is Perceptron?
  - How is it defined?
  - Build deep NN from a Perceptron



# The Perceptron – A single Neuron



Linear combination of inputs

Output

$\hat{y} = g(w_0 + \sum_{i=1}^m x_i w_i)$

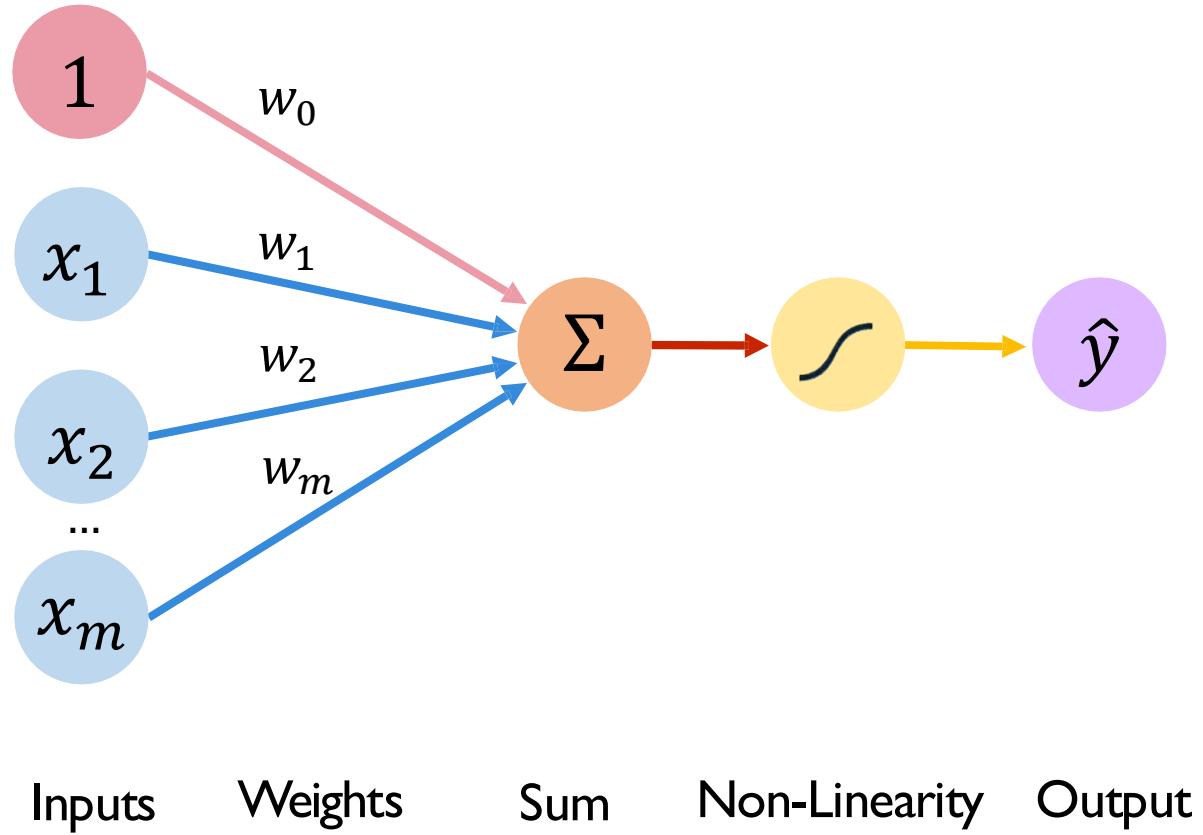
Non-linear activation function

Bias

Binary Classification:

$(x_1, \dots, x_m)$  feature of a data  
 $\hat{y} \in [0,1]$ : probability belongs to the positive class

# The Perceptron

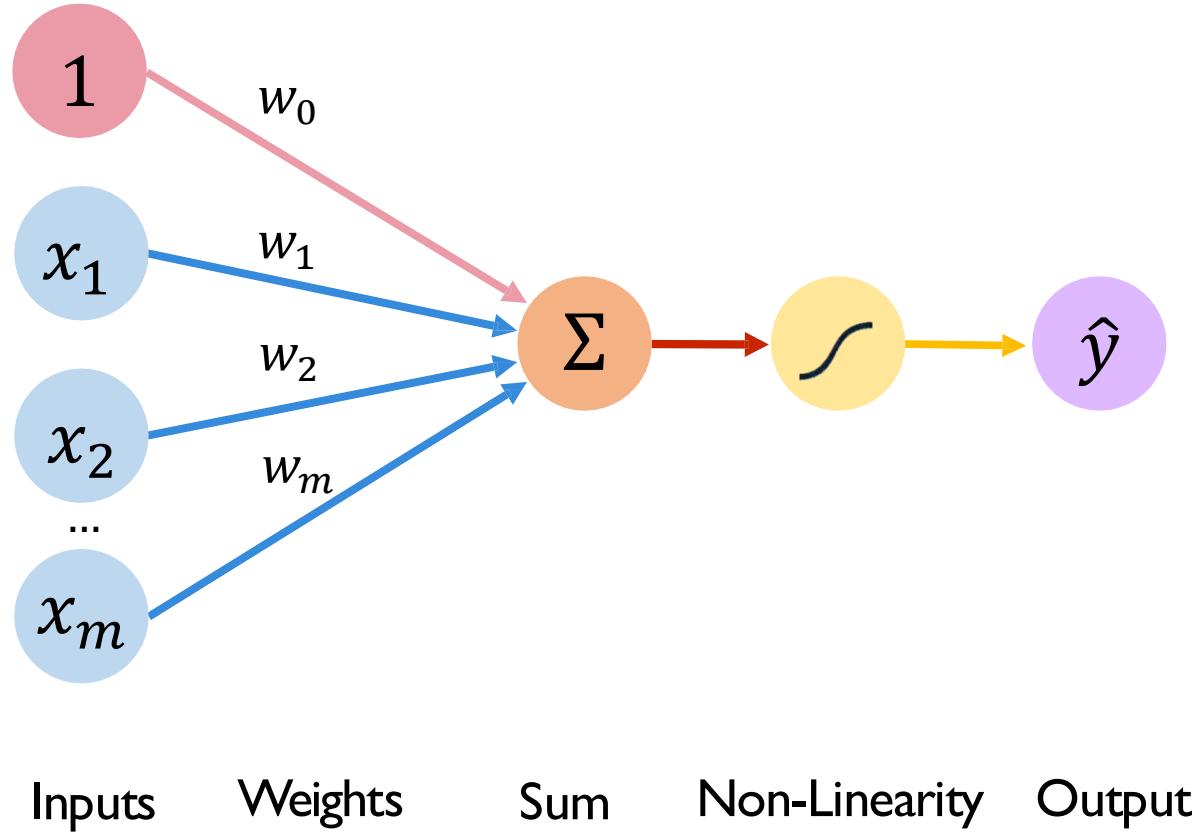


$$\hat{y} = g(w_0 + \sum_{i=1}^m x_i w_i)$$

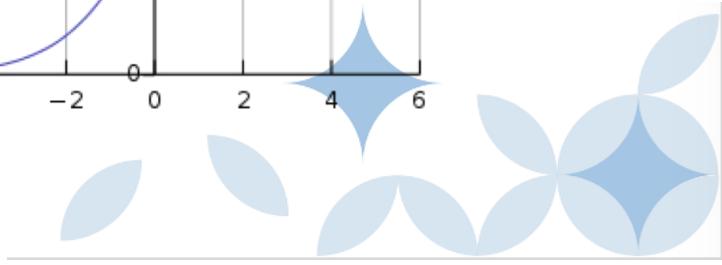
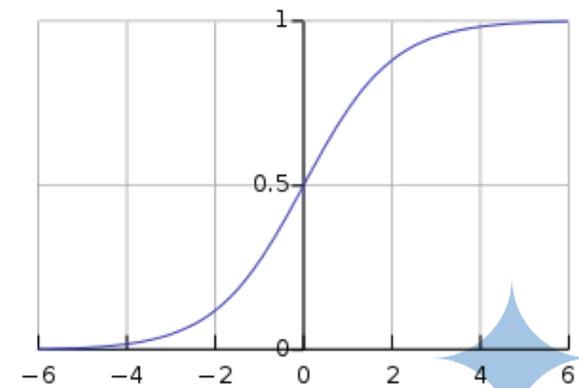
$$\hat{y} = g(w_0 + X^T W)$$

where:  $X = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$  and  $W = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

# The Perceptron



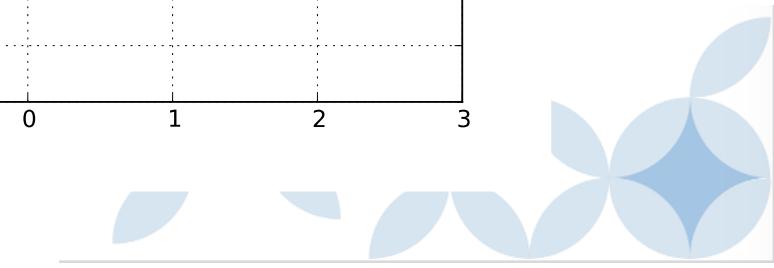
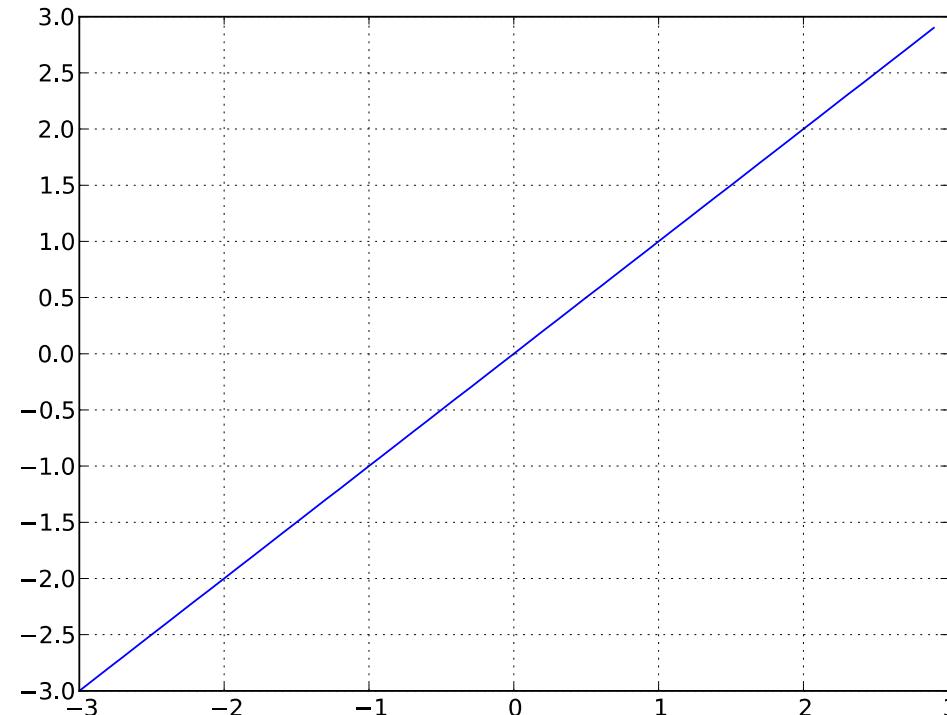
$$\hat{y} = g(w_0 + \sum_{i=1}^m x_i w_i)$$



# Activation Functions

Linear activation function:

- $g(a) = a$
- Correspond to linear regression



# Non-linear Activation Functions

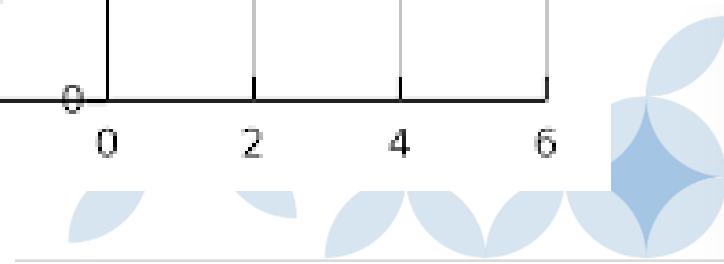
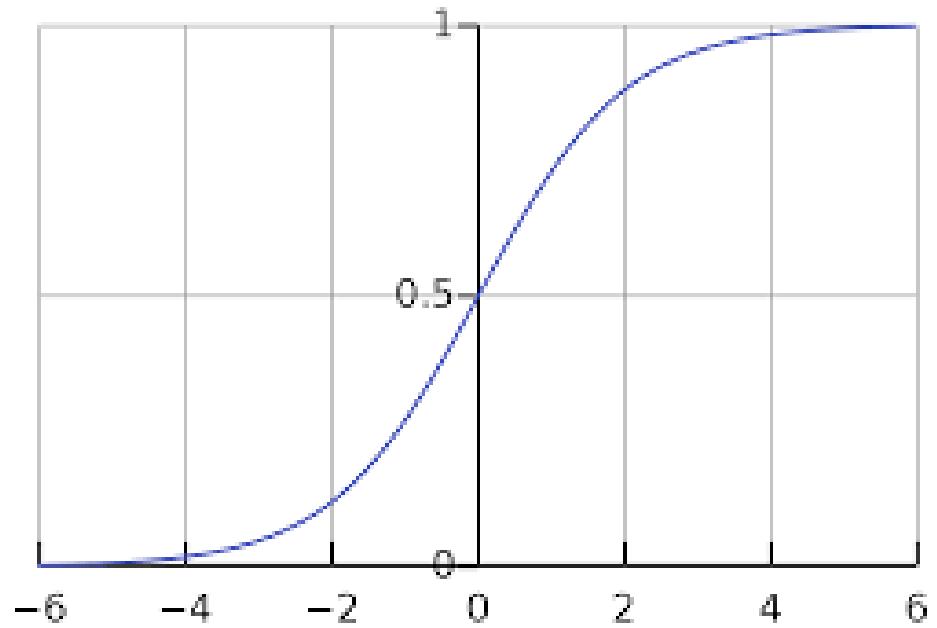
## Sigmoid activation function:

- Squashes the neuron's output between 0 and 1

$$g(a) = \sigma(a) = \frac{\exp(a)}{1 + \exp(a)} = \frac{1}{1 + \exp(-a)}$$

- Always positive and strictly increasing
- Naturally suitable for probability
  - Logistic Regression

$$\hat{y} = P(Y = 1|X) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

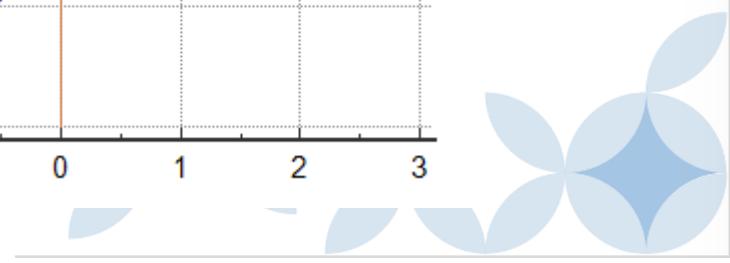
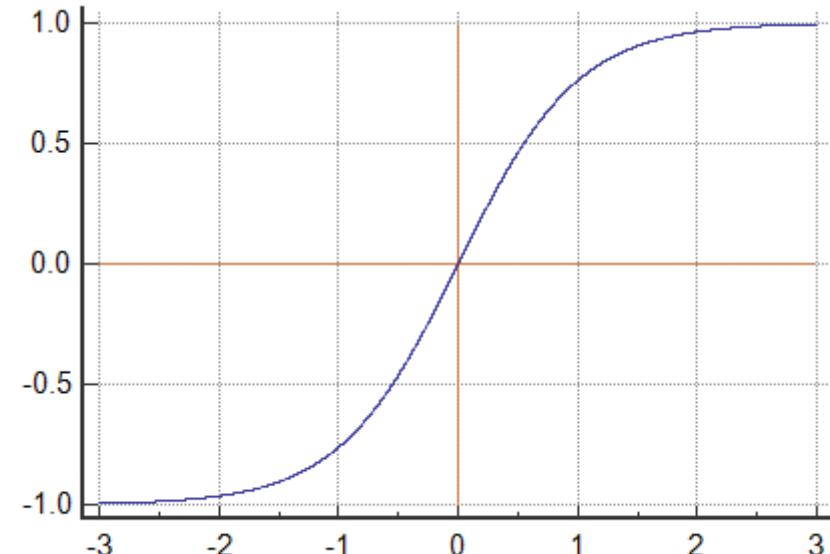


# Non-linear Activation Functions

Hyperbolic tangent (“tanh”) activation function:

$$g(a) = \tanh(a) = \frac{\exp(a) - \exp(-a)}{\exp(a) + \exp(-a)} = \frac{\exp(2a) - 1}{\exp(2a) + 1}$$

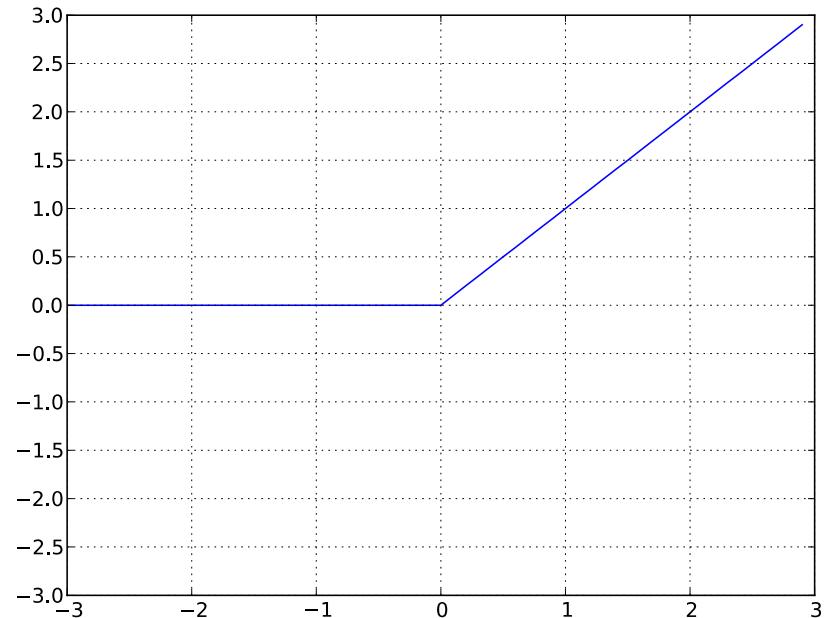
- Squashes the neuron’s output between -1 and 1
- Strictly Increasing



# Non-linear Activation Functions

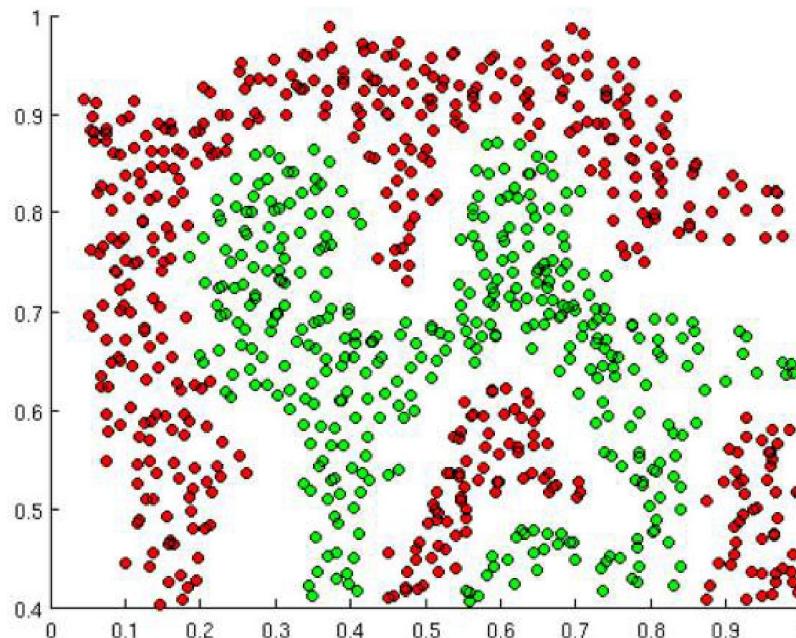
## Rectified linear Unit (ReLU), activation function:

- Rectified linear Unit (ReLU), activation function:
  - Nair and Hinton (2010)
  - Bounded below by 0 (always non-negative)
  - Not upper bounded
  - Not smooth
  - other variants of ReLU ( Leaky ReLU )



# Importance of Activation Function

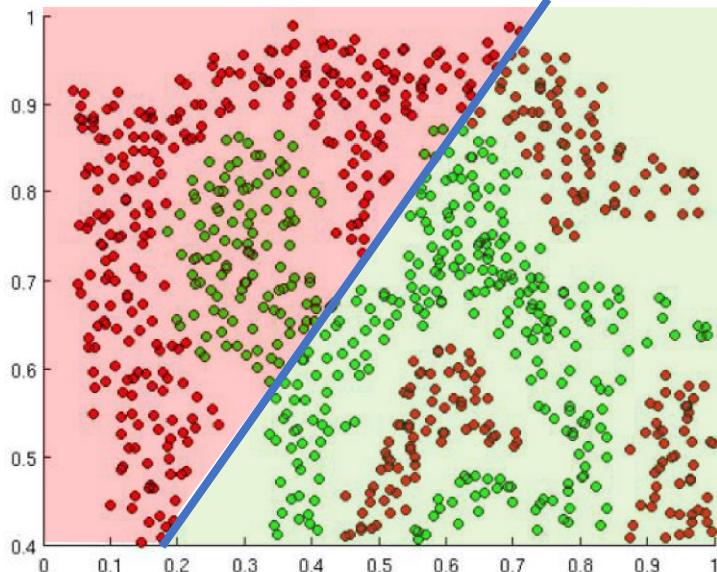
## Non-linearity



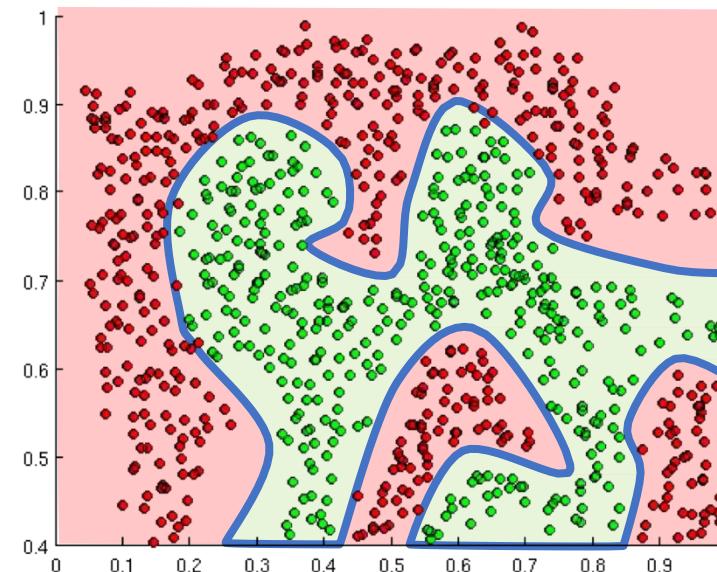
We want to build a Neural Network to distinguish green vs red points



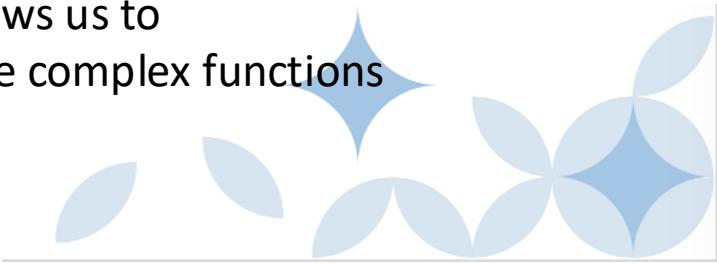
# More complex neural networks



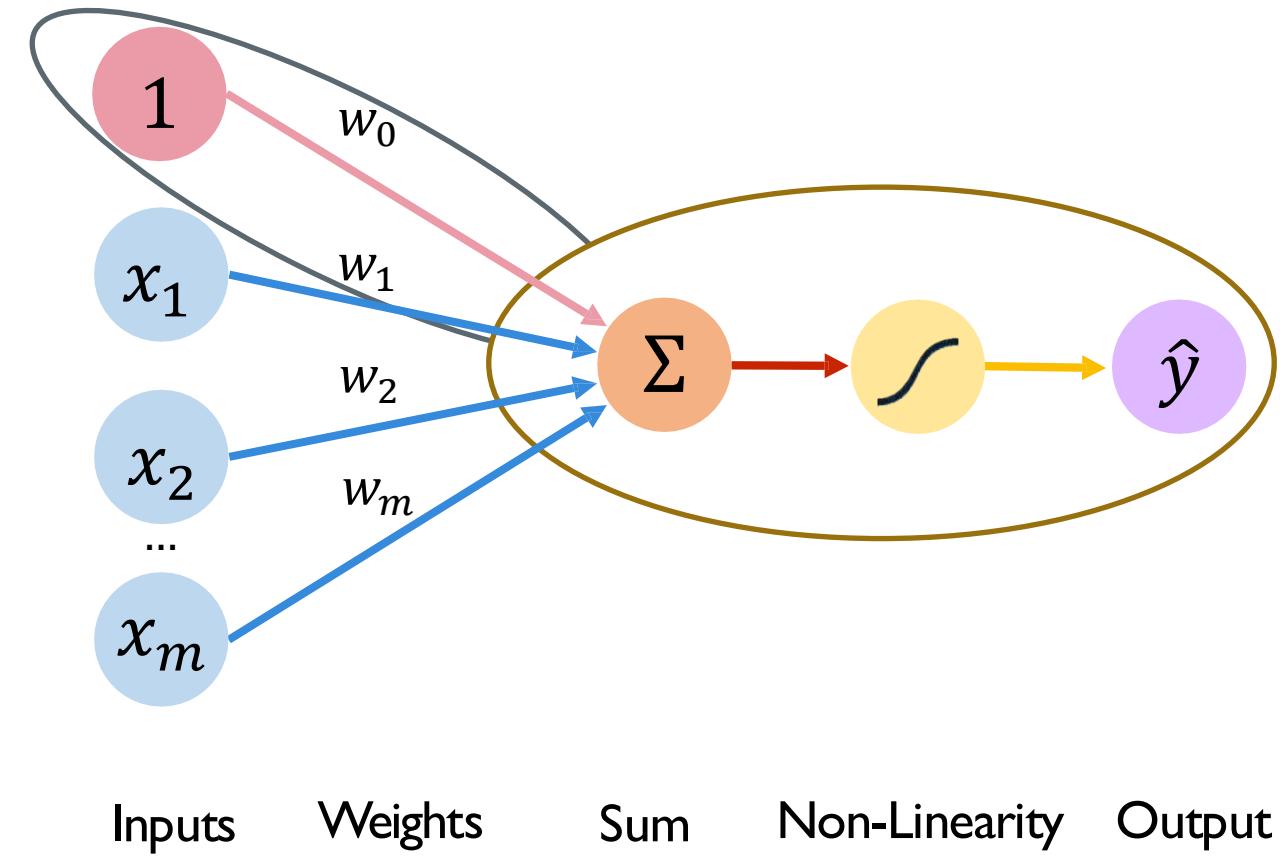
Linear activation function produce  
linear decision boundaries



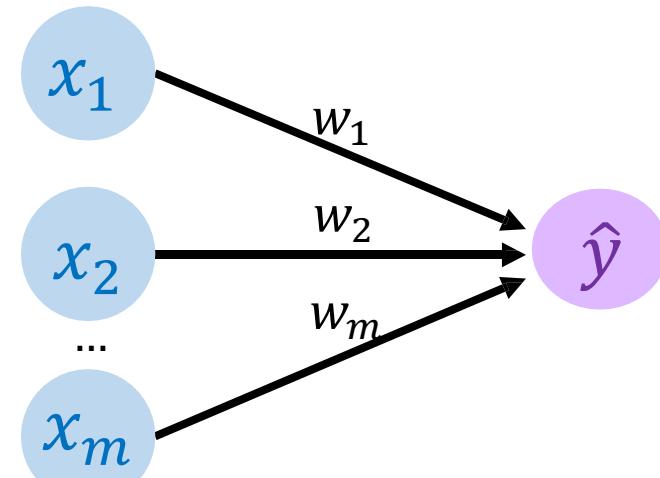
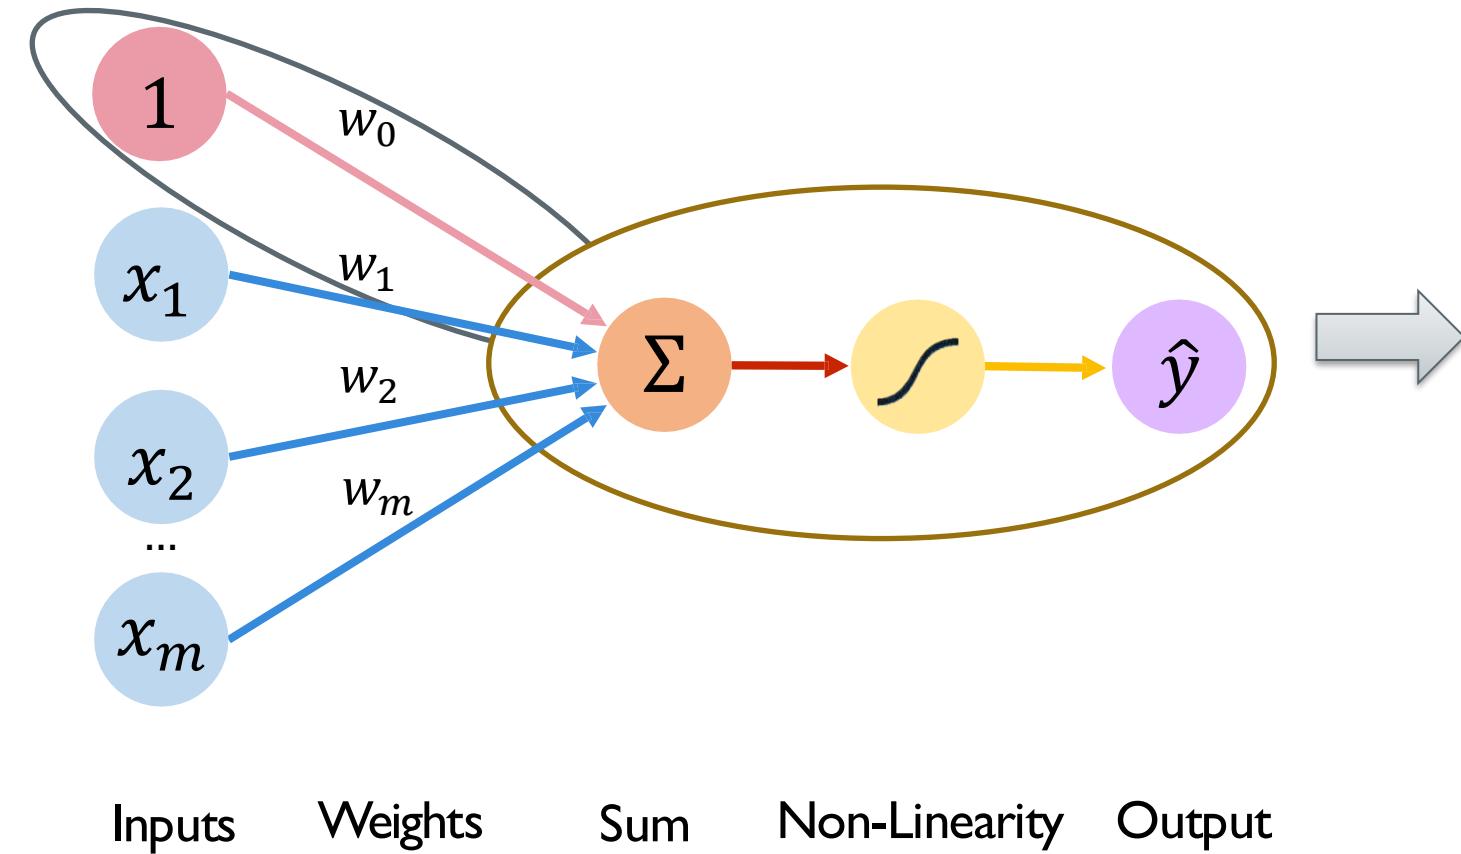
Non-linearity allows us to  
approximate complex functions



# The Perceptron:Simplified

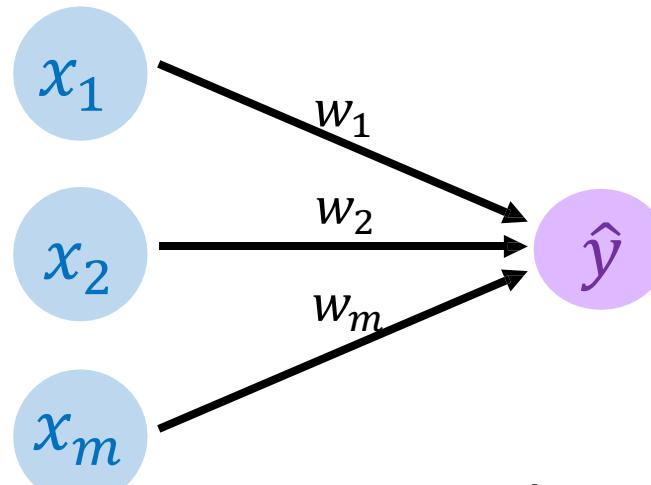


# The Perceptron: Simplified



$$\hat{y} = g(w_0 + \sum_{i=1}^m x_i w_i)$$

# Single Output Perceptron

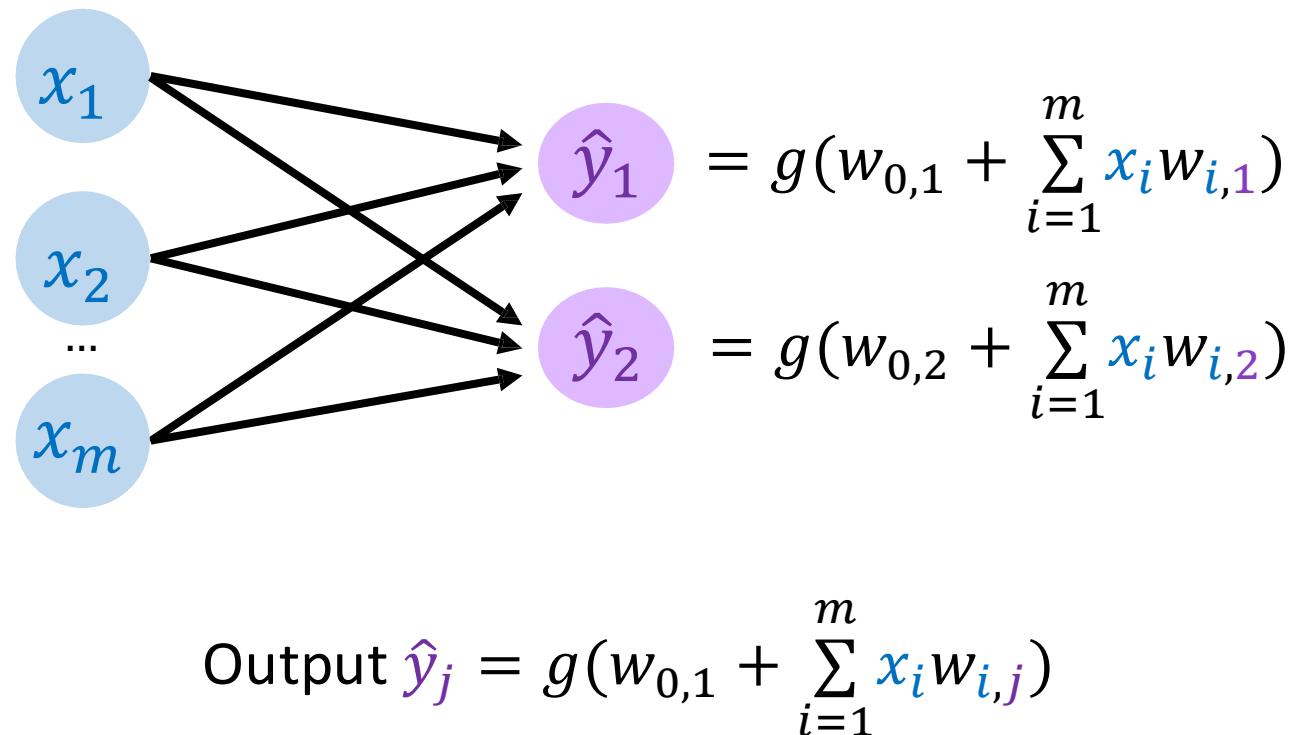


$$\hat{y} = g(w_0 + \sum_{i=1}^m x_i w_i)$$

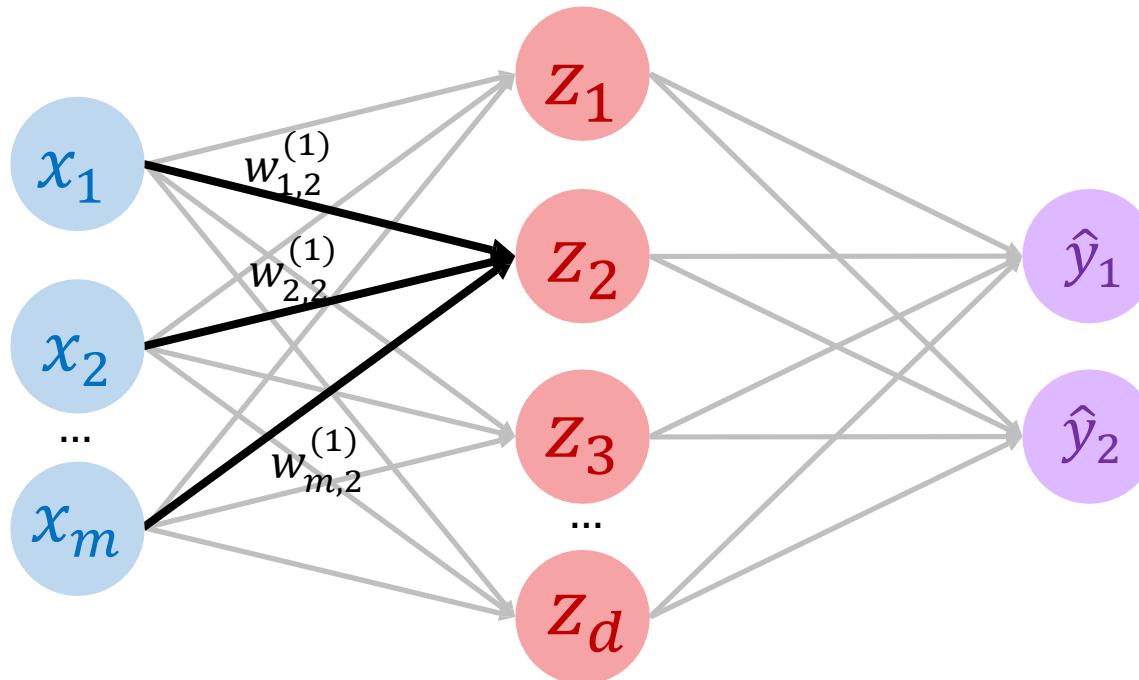
What if we have multi-output  $\hat{y}_1$  and  $\hat{y}_2$ ?



# Multi Output Perceptron



# Extend to Single (Hidden) Layer Neural Network



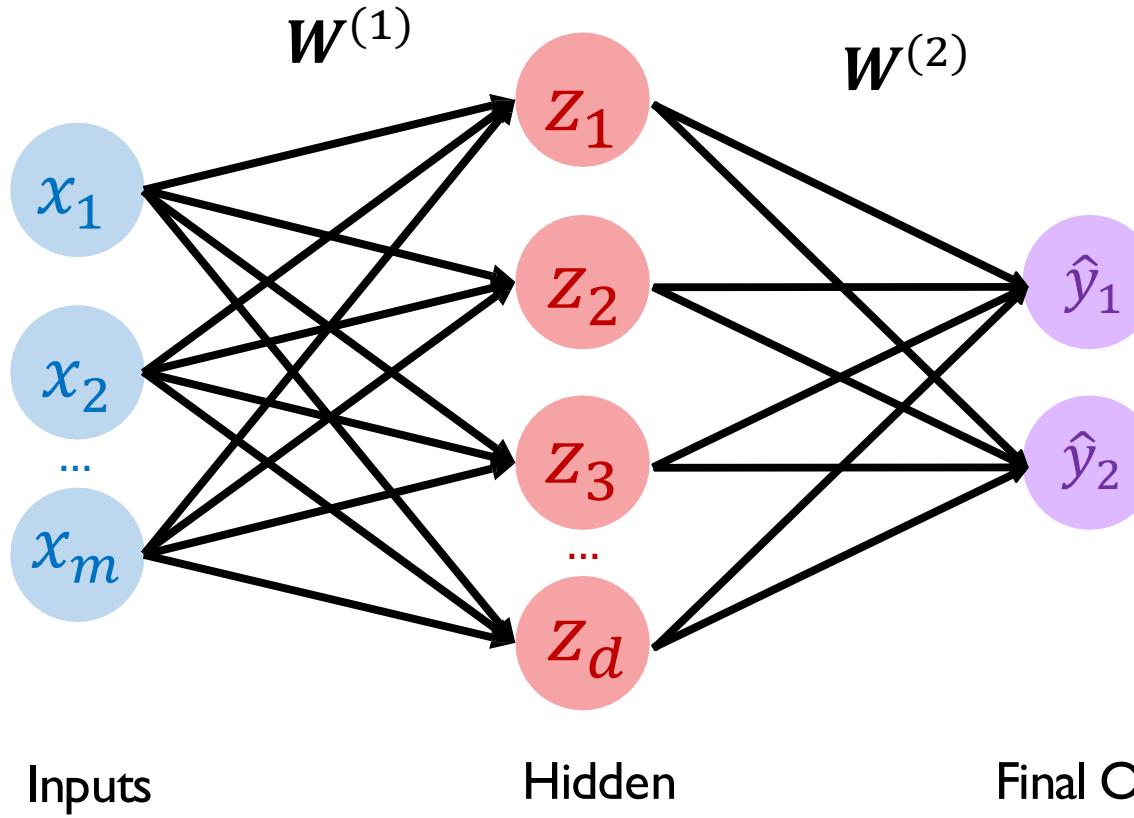
$$z_2 = g(w_{0,2}^{(1)} + \sum_{i=1}^m x_i w_{i,2}^{(1)})$$

$$= g(w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + \cdots + x_m w_{m,2}^{(1)})$$

(1) Indicates the 1st layer



# Single (Hidden) Layer Neural Network



Parameters to be learned

$$\mathbf{W} = (\mathbf{W}^{(1)}, \mathbf{W}^{(2)})$$

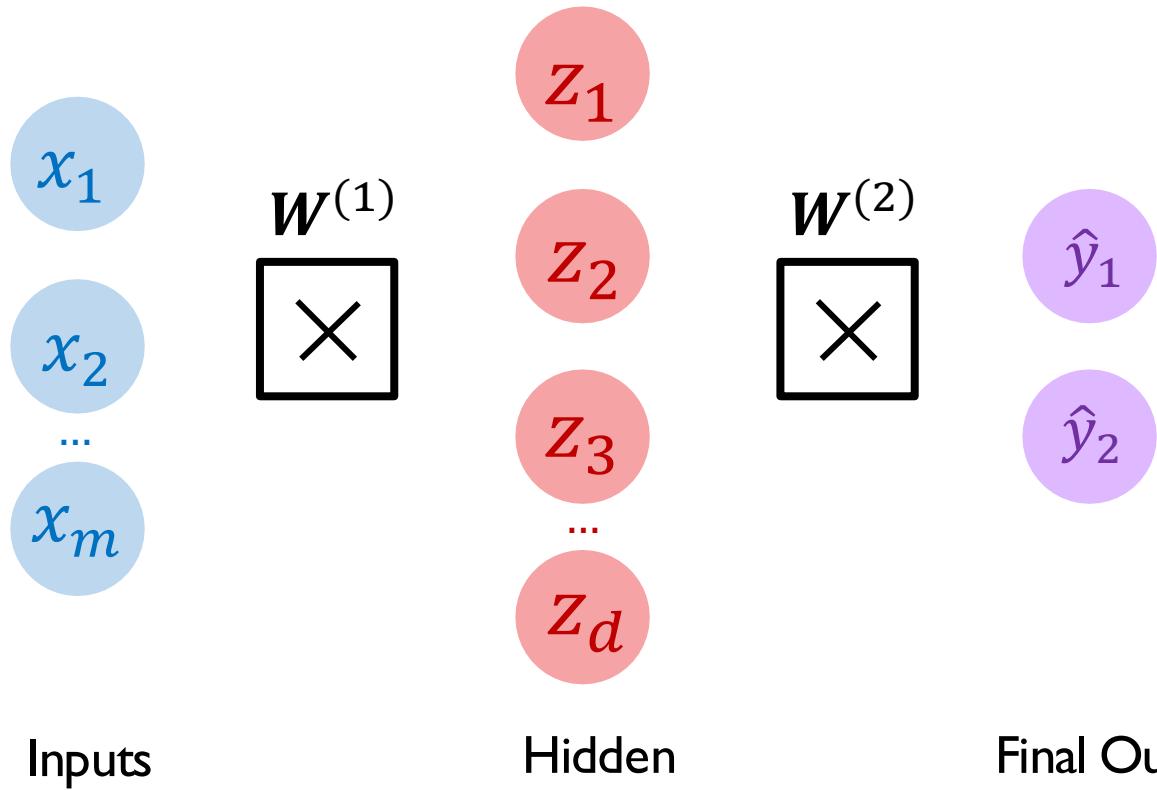
$$z_j = g(w_{0,j}^{(1)} + \sum_{i=1}^m x_i w_{i,j}^{(1)})$$

Hidden

$$\hat{y}_k = g(w_{0,k}^{(2)} + \sum_{j=1}^d z_j w_{j,k}^{(2)})$$



# Single (Hidden) Layer Neural Network



Inputs

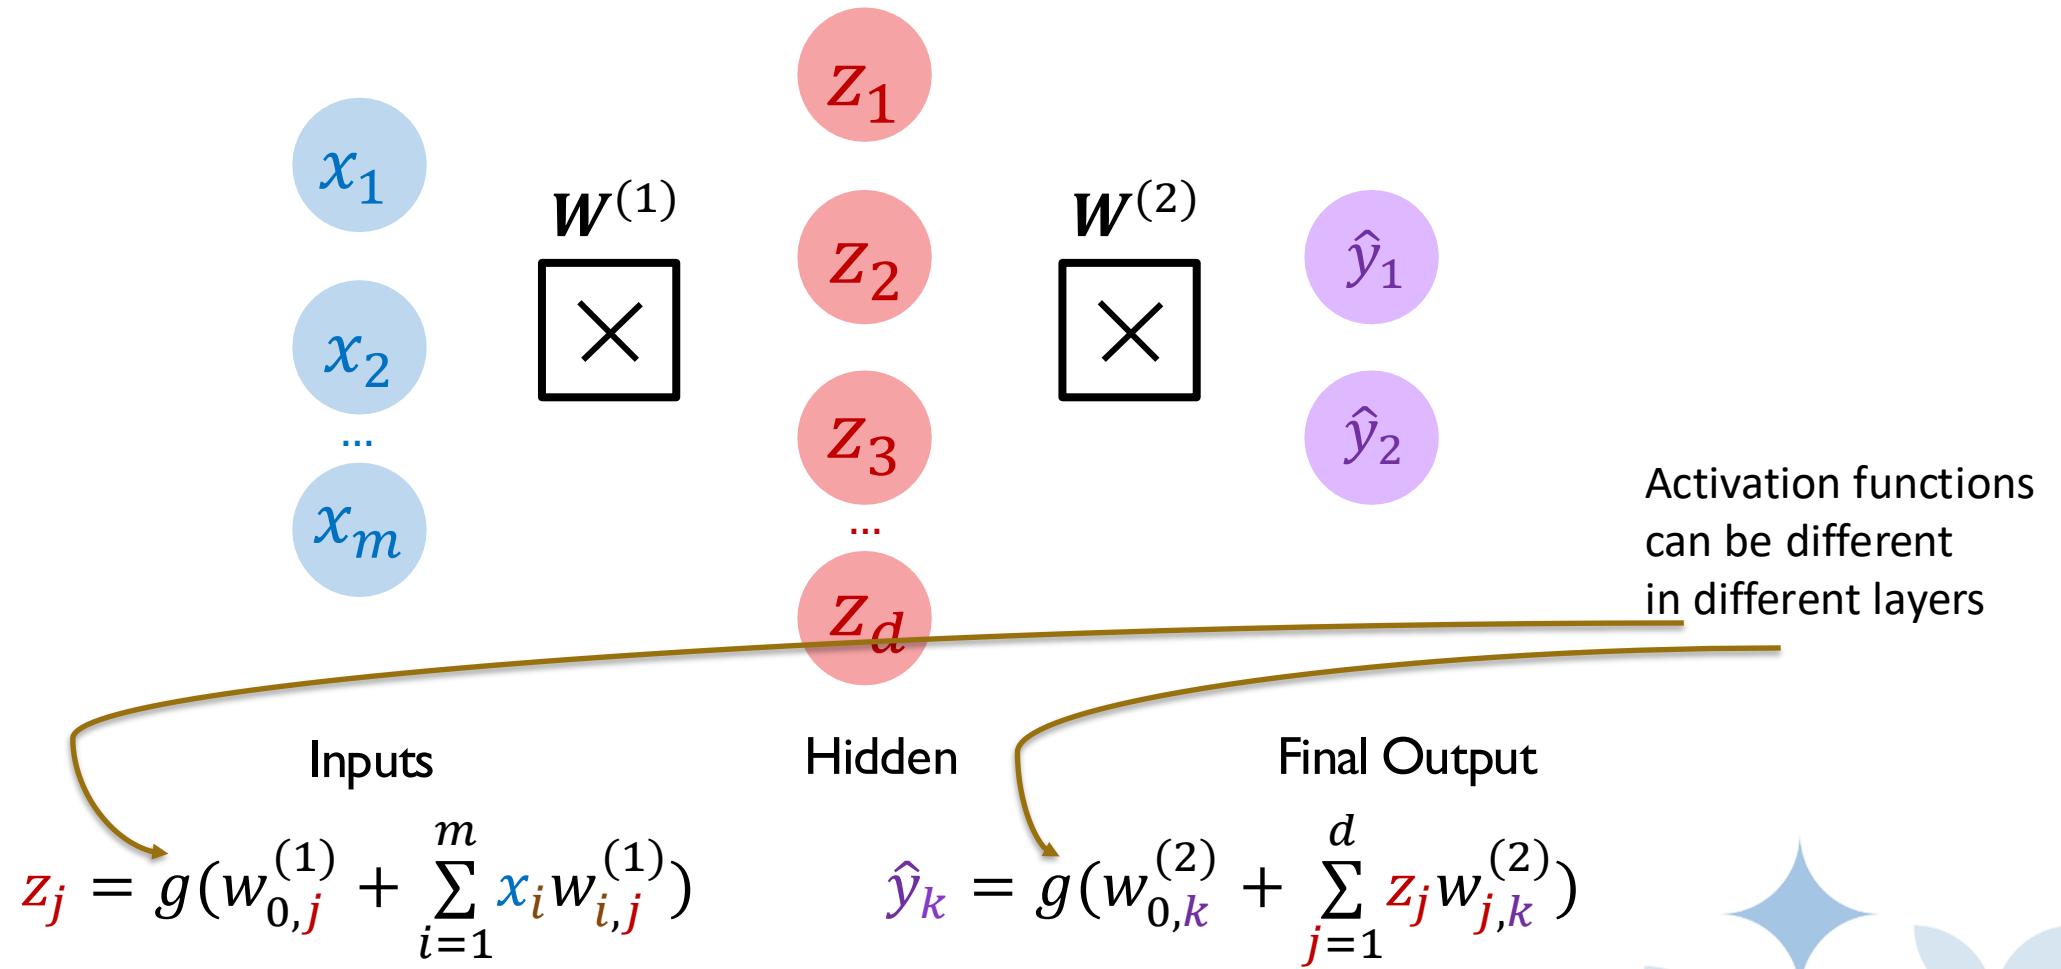
$$z_j = g(w_{0,j}^{(1)} + \sum_{i=1}^m x_i w_{i,j}^{(1)})$$

Hidden

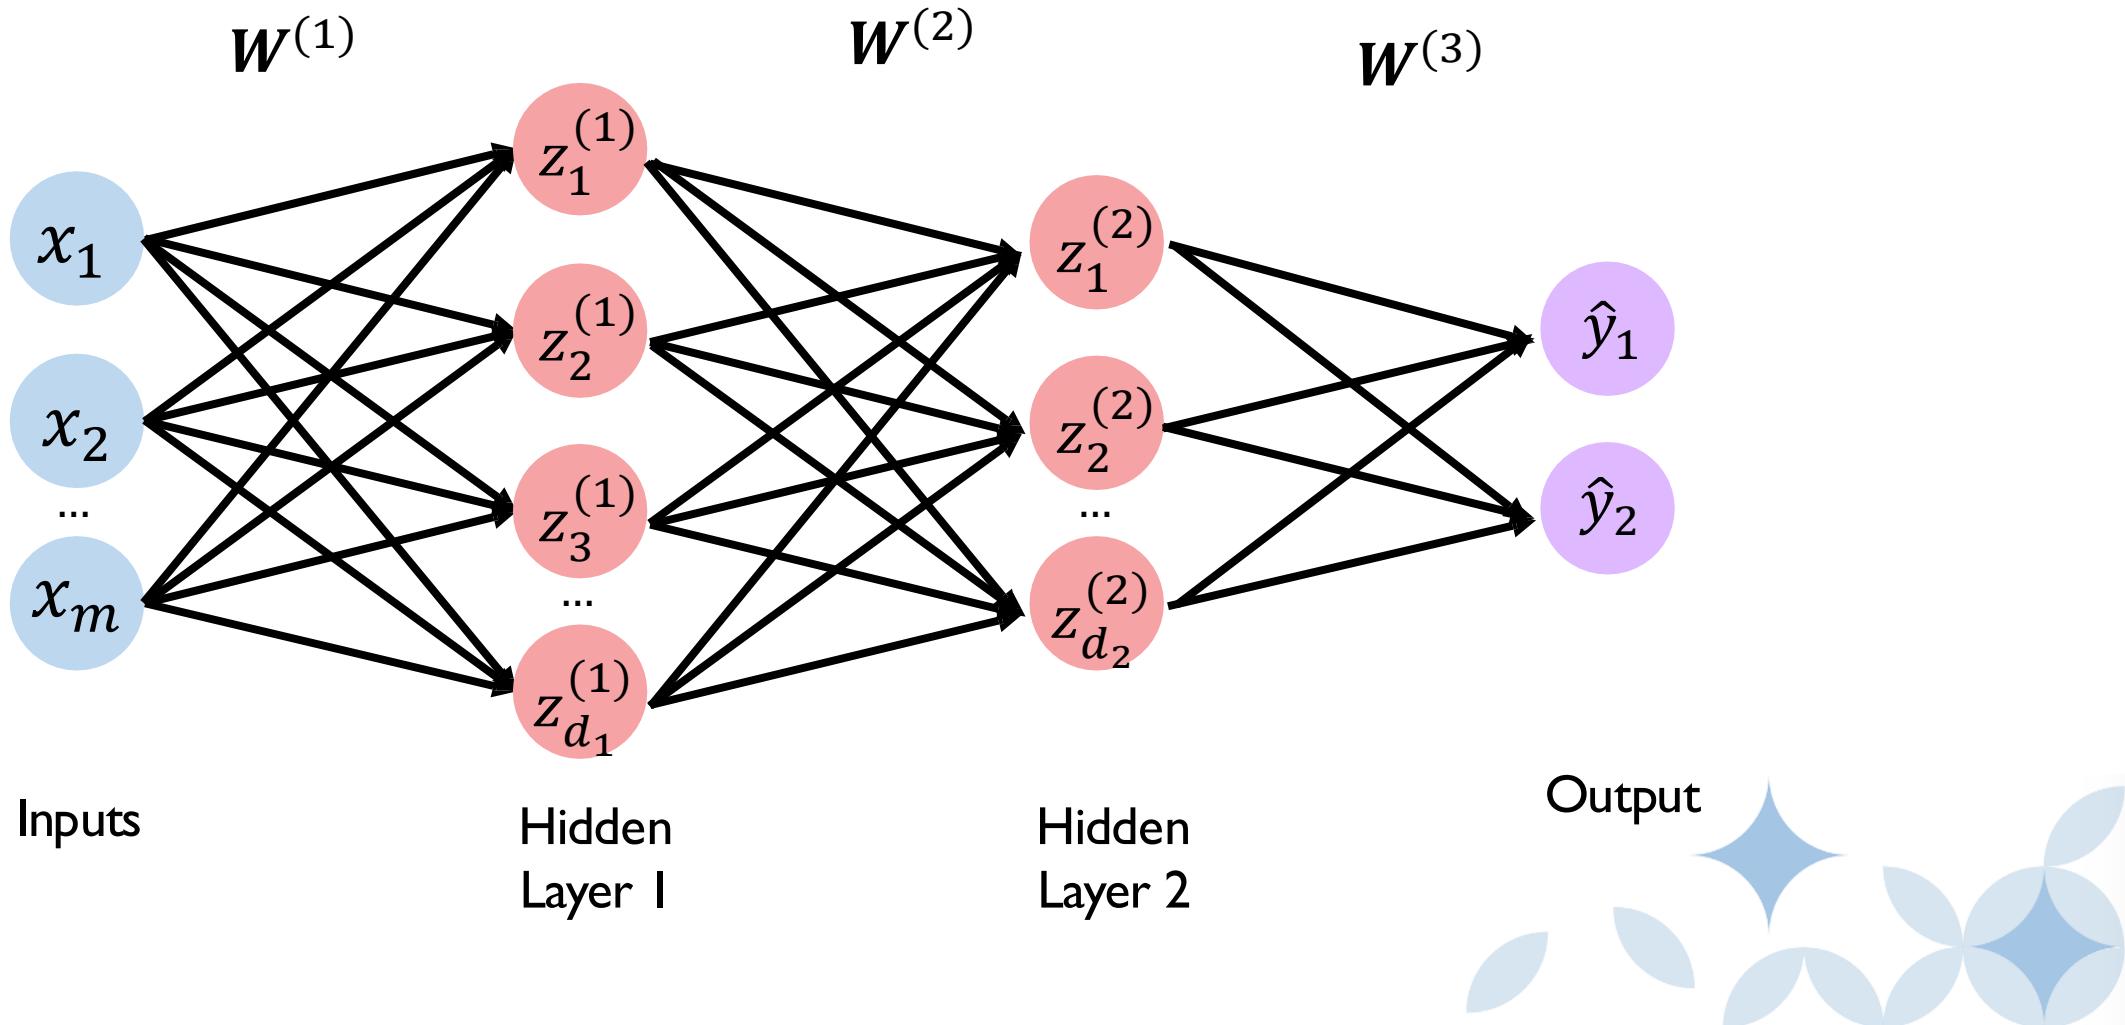
$$\hat{y}_k = g(w_{0,k}^{(2)} + \sum_{j=1}^d z_j w_{j,k}^{(2)})$$



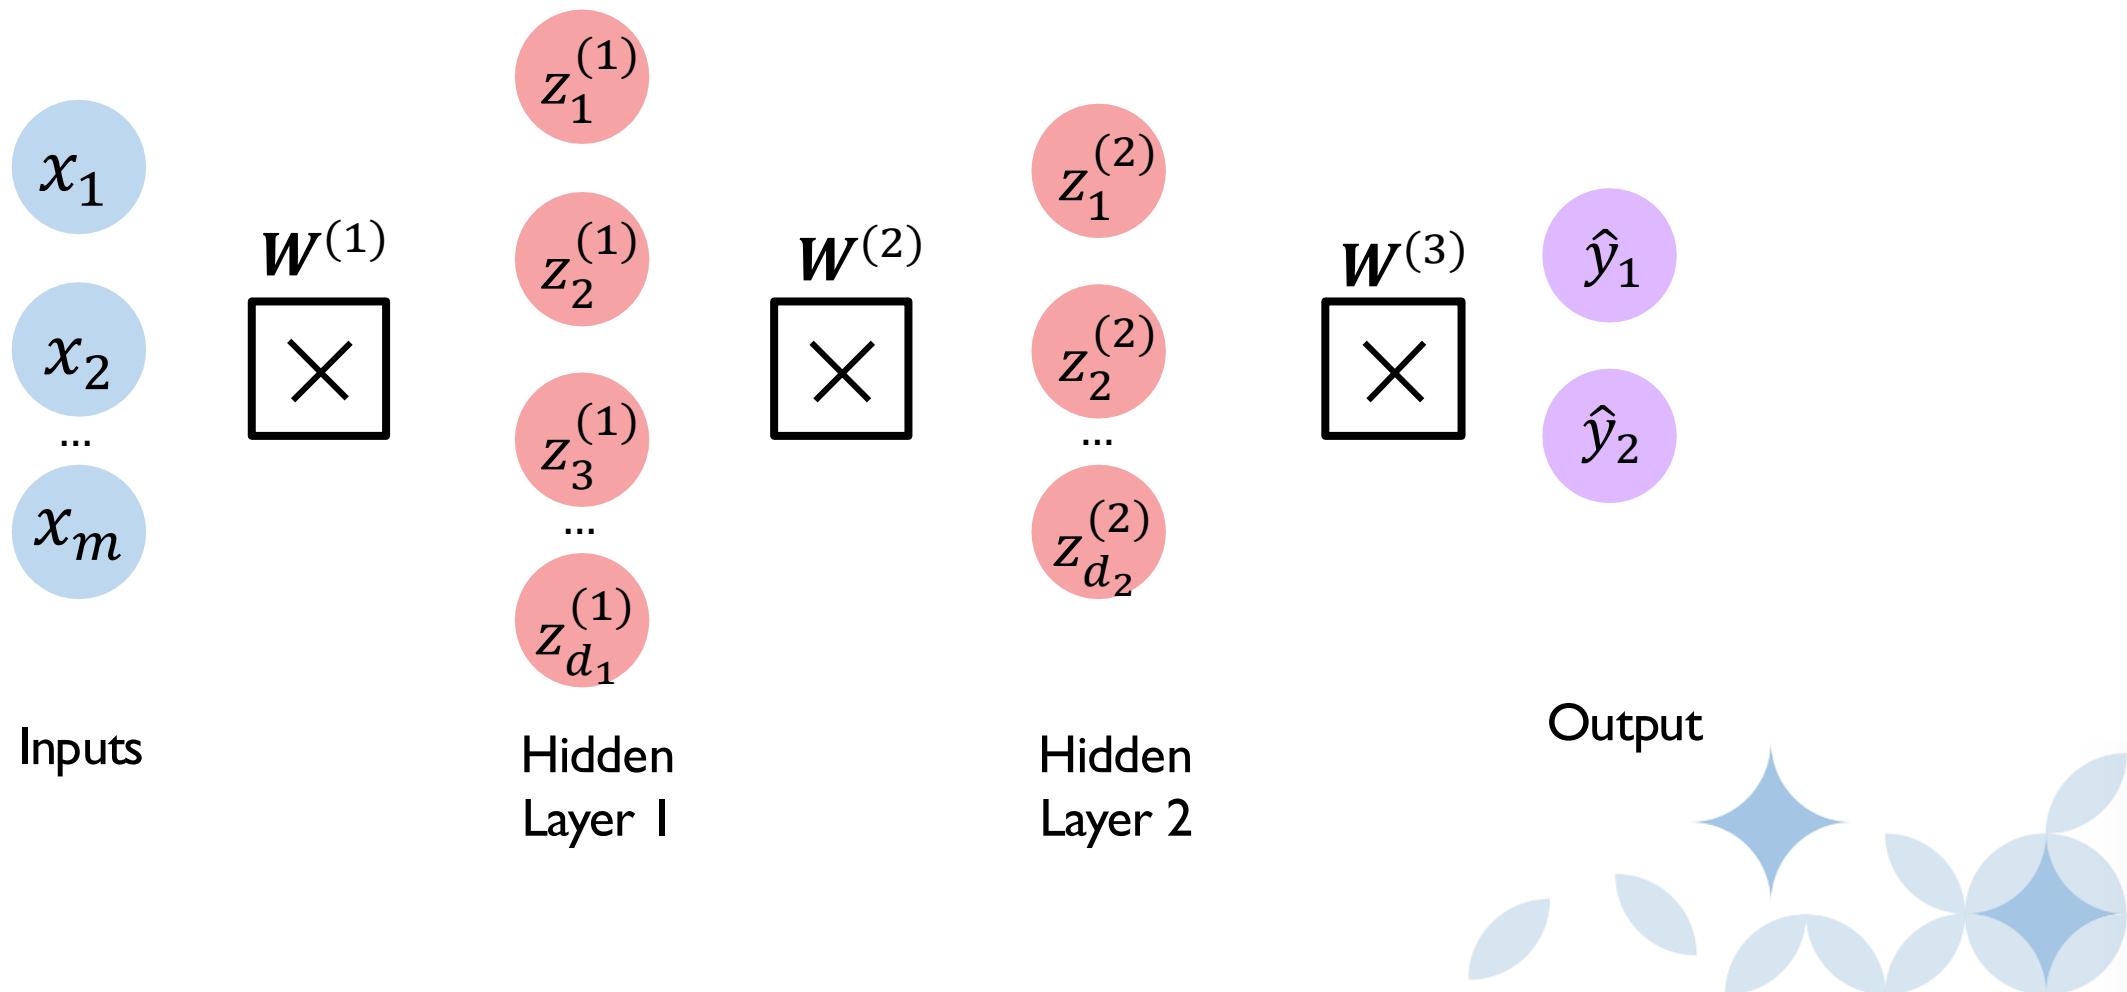
# Single (Hidden) Layer Neural Network



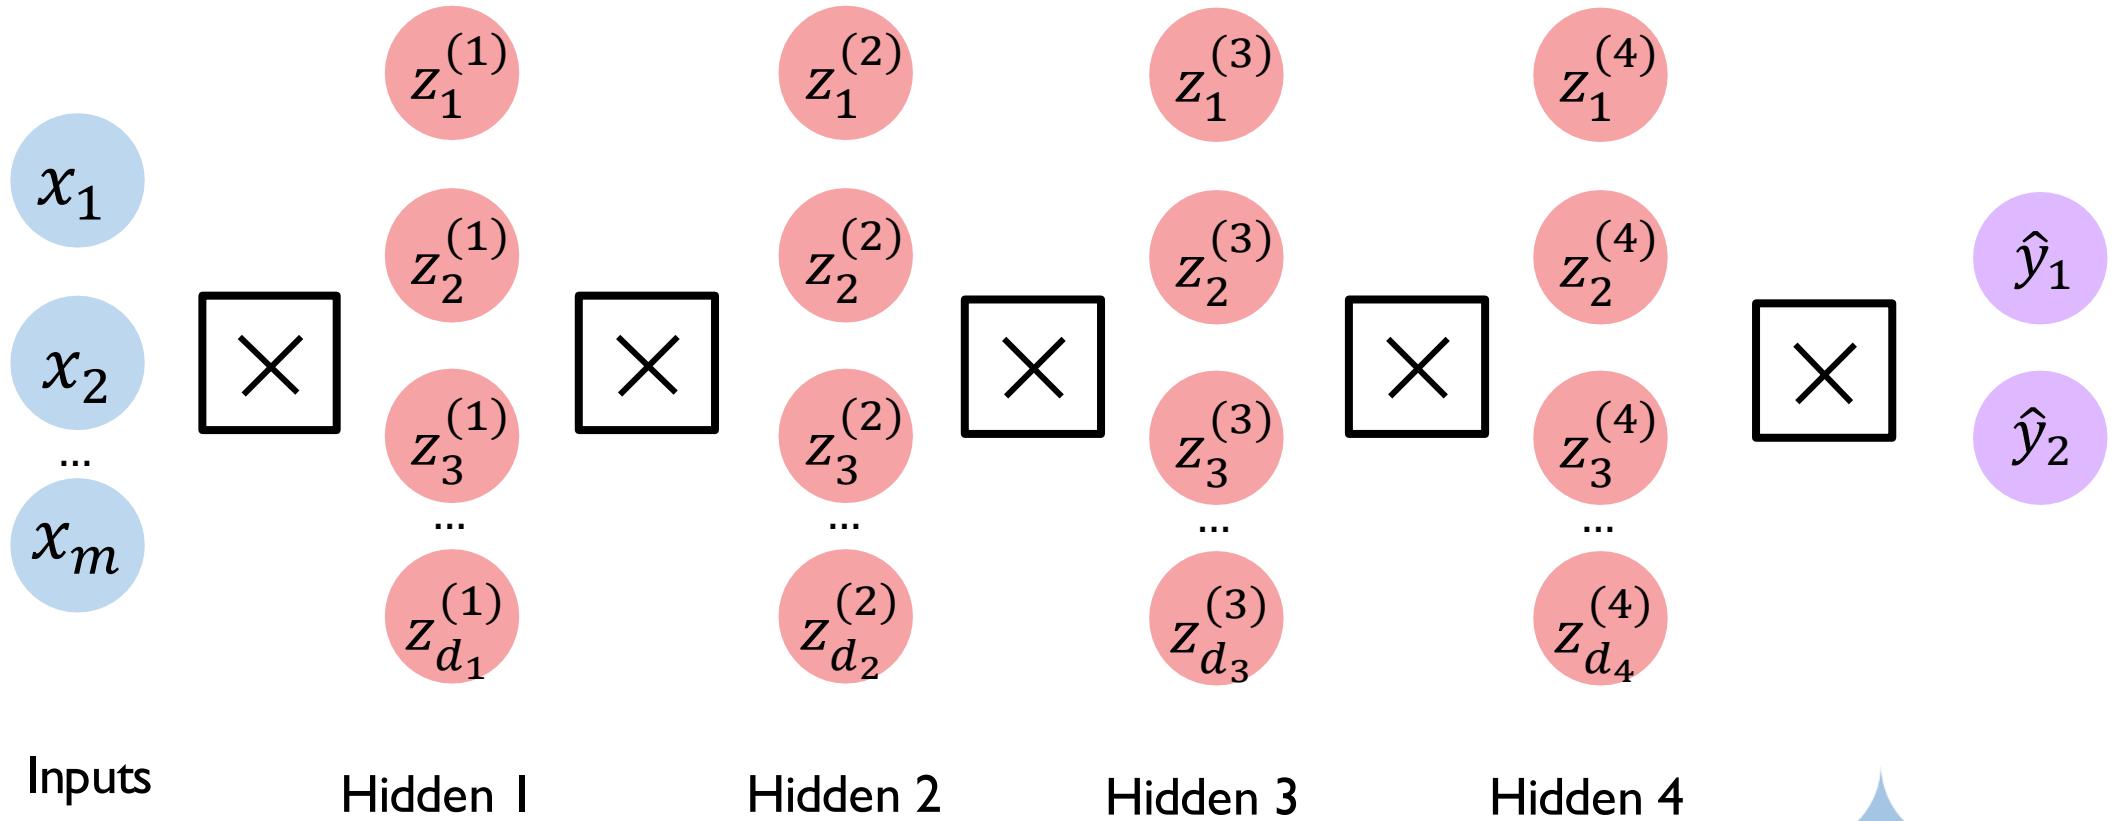
# Two-Layer Neural Network



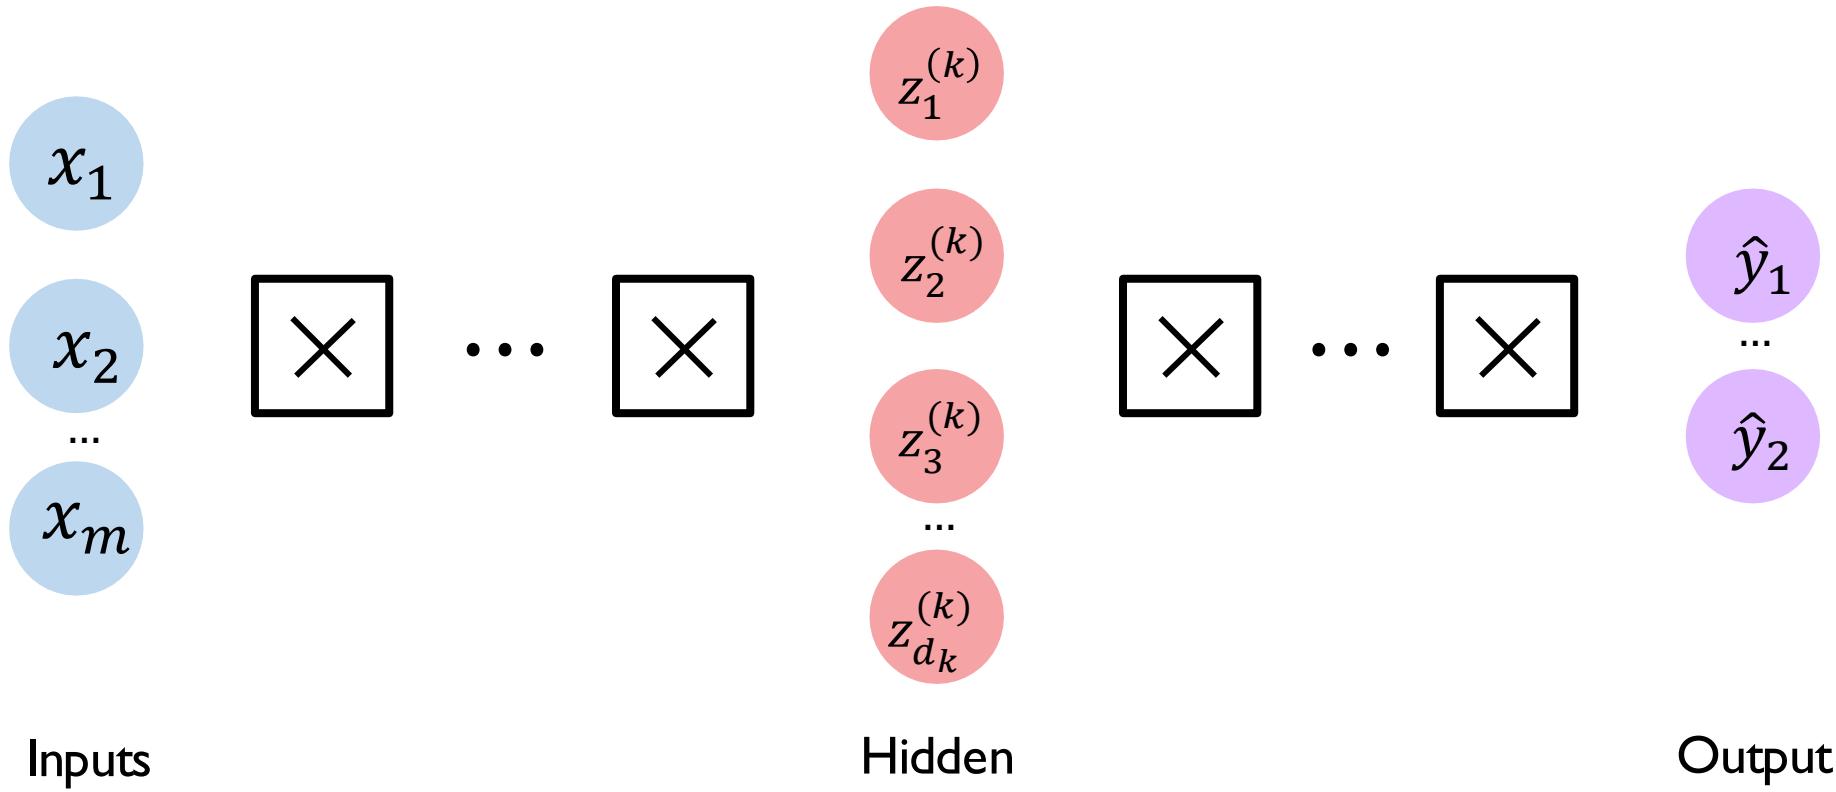
# Two-Layer Neural Network



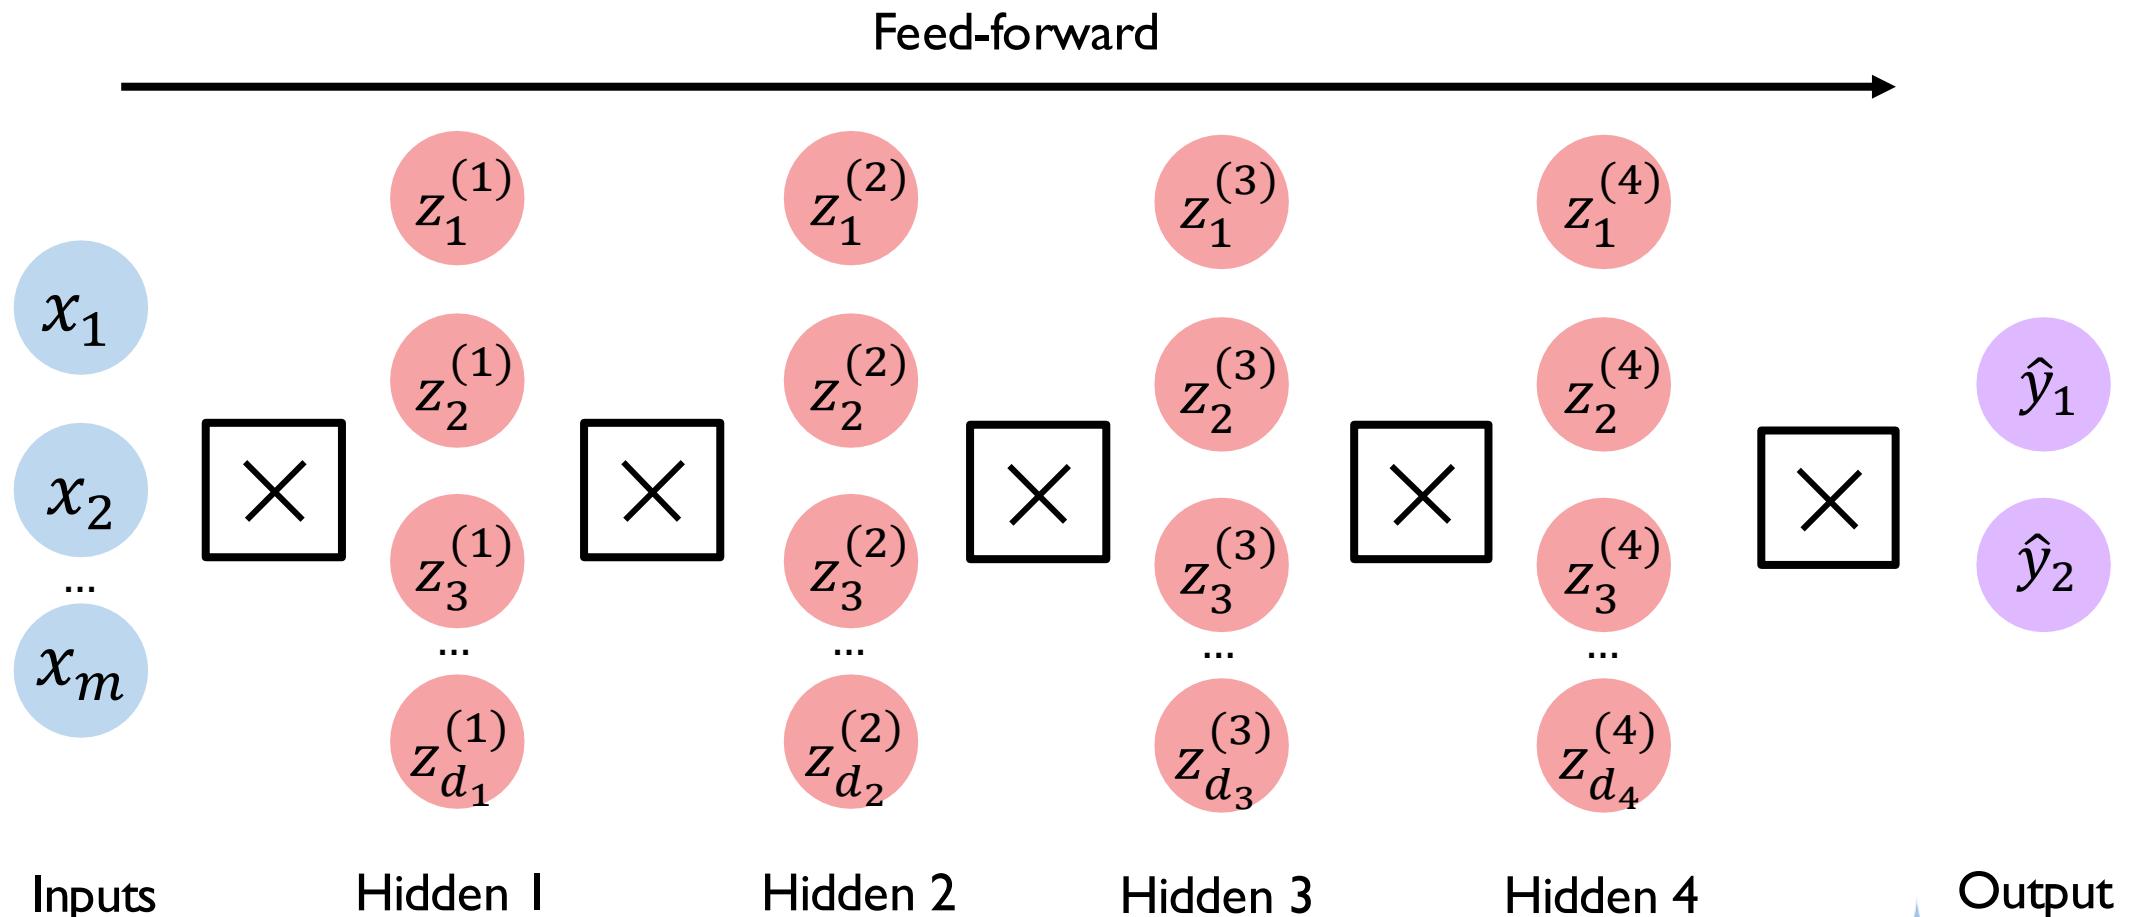
# Multi-Layer (Deep) Neural Network



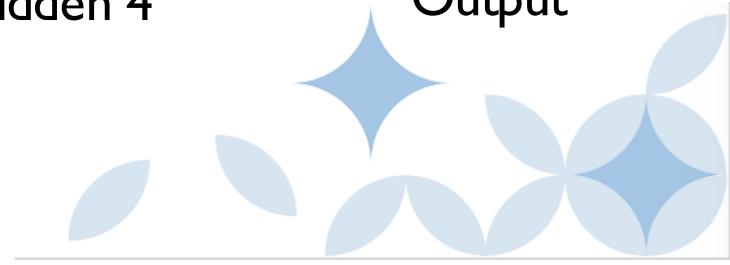
# Deep Neural Network (DNN)



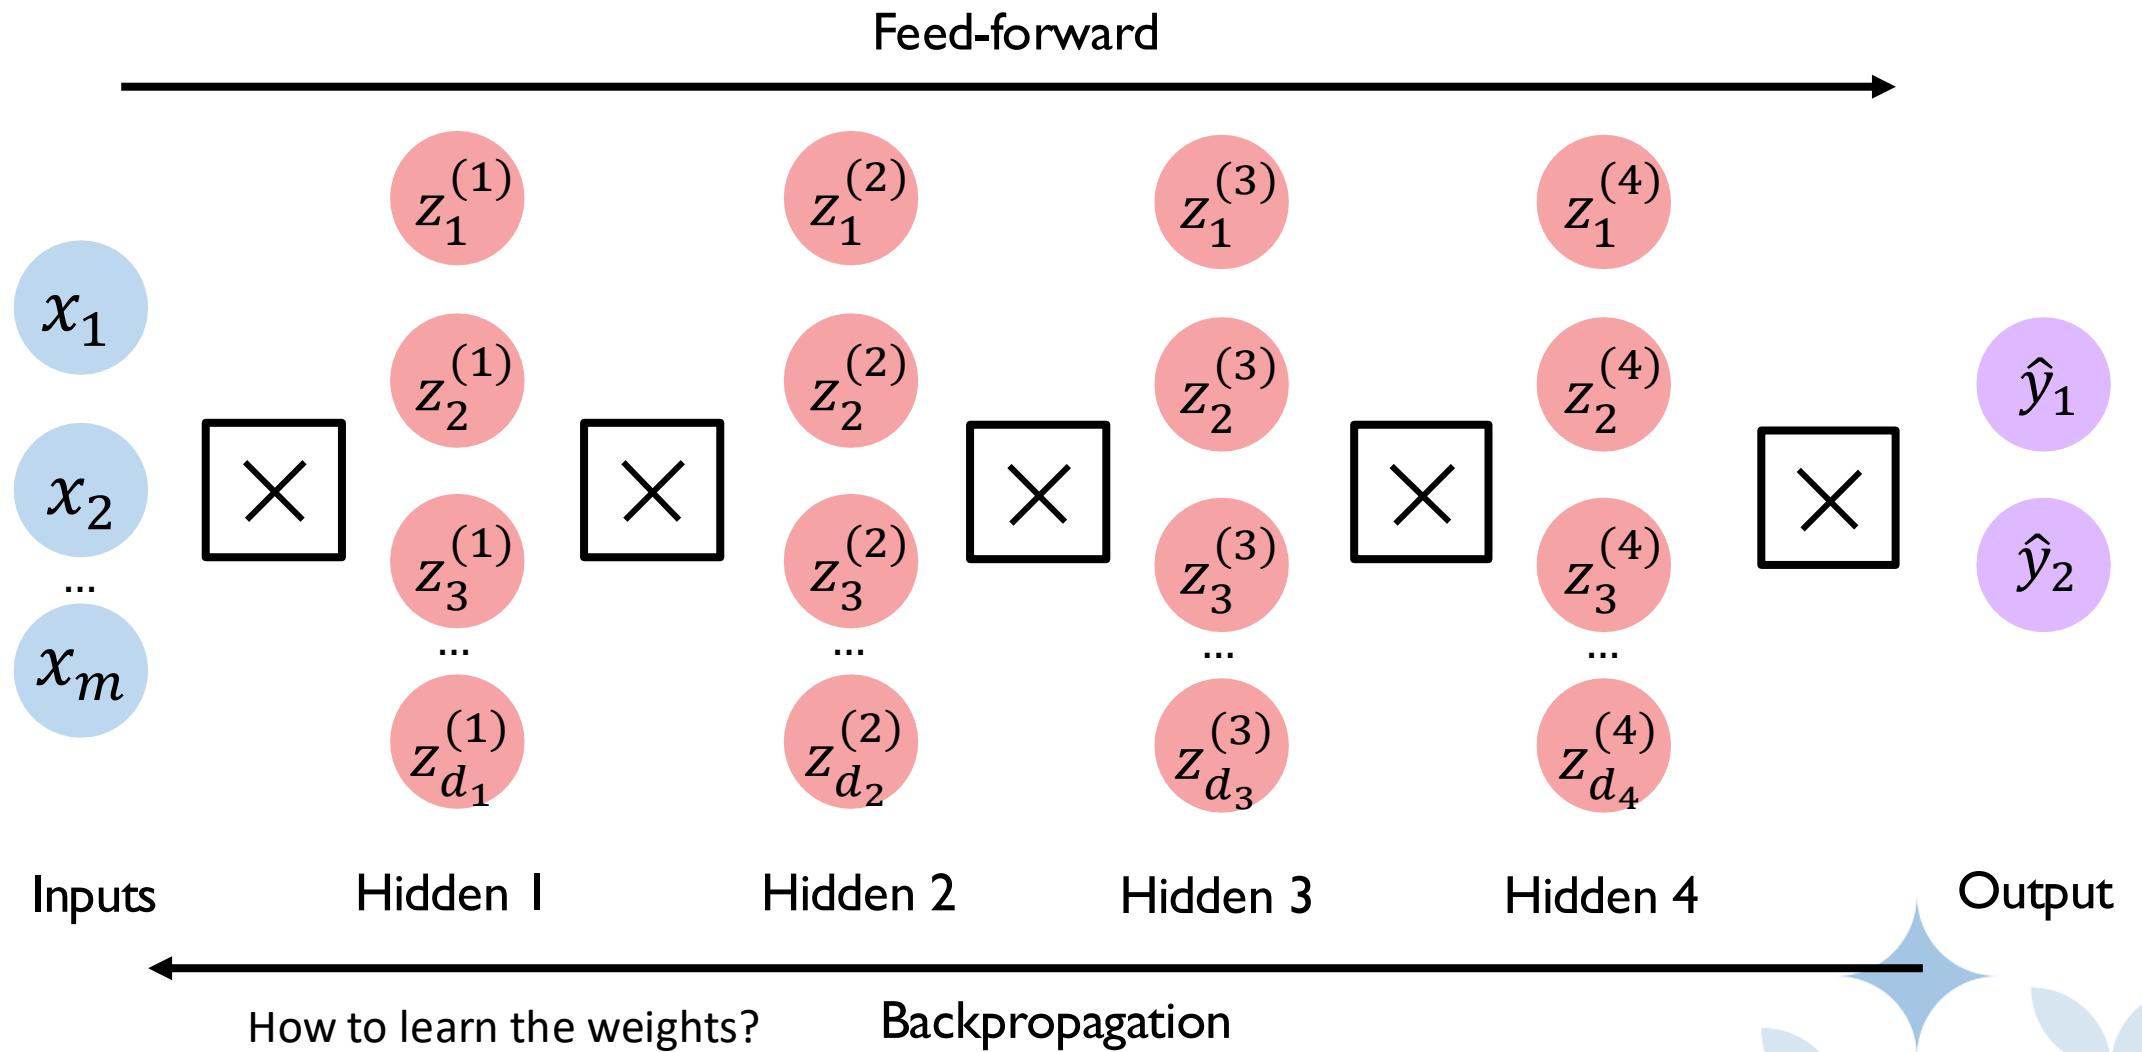
# Feed-forward Networks and Backpropagation



How to learn the weights?



# Feed-forward Networks and Backpropagation



# Example Problem

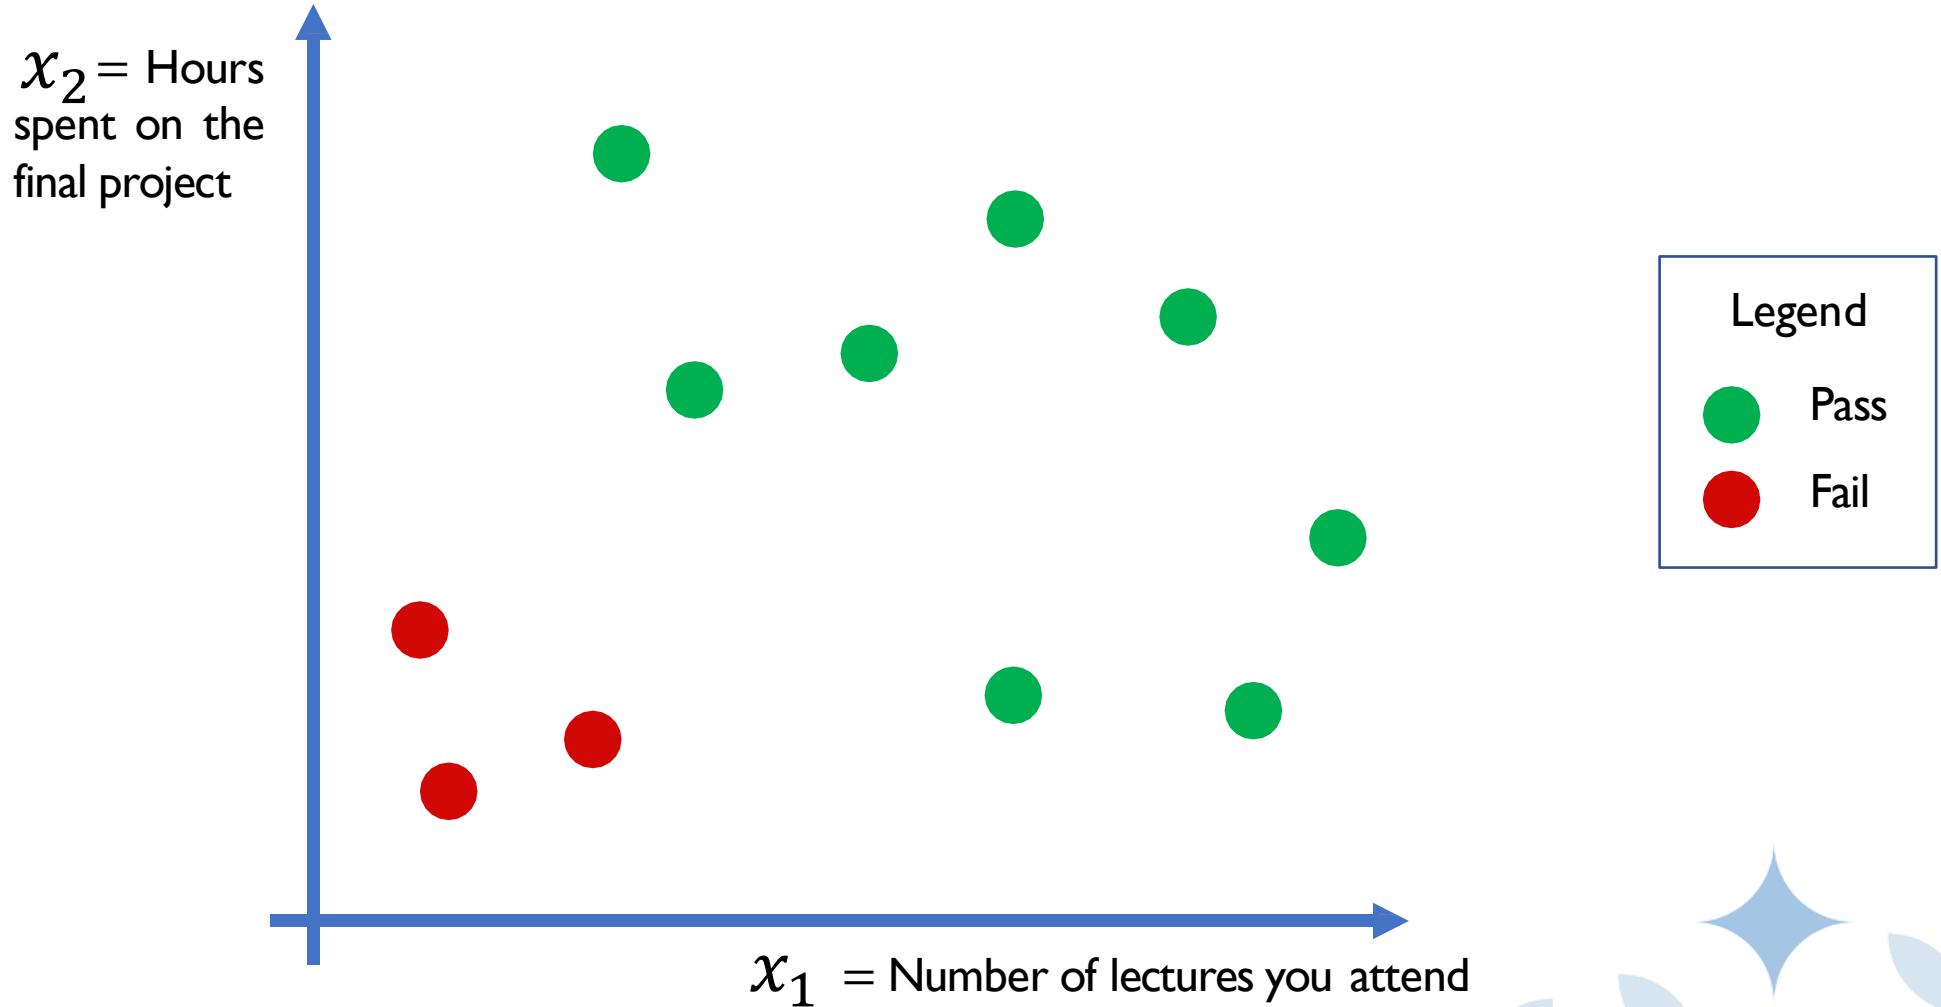
## Will I pass the class?

– Let's start with a simple two-feature model

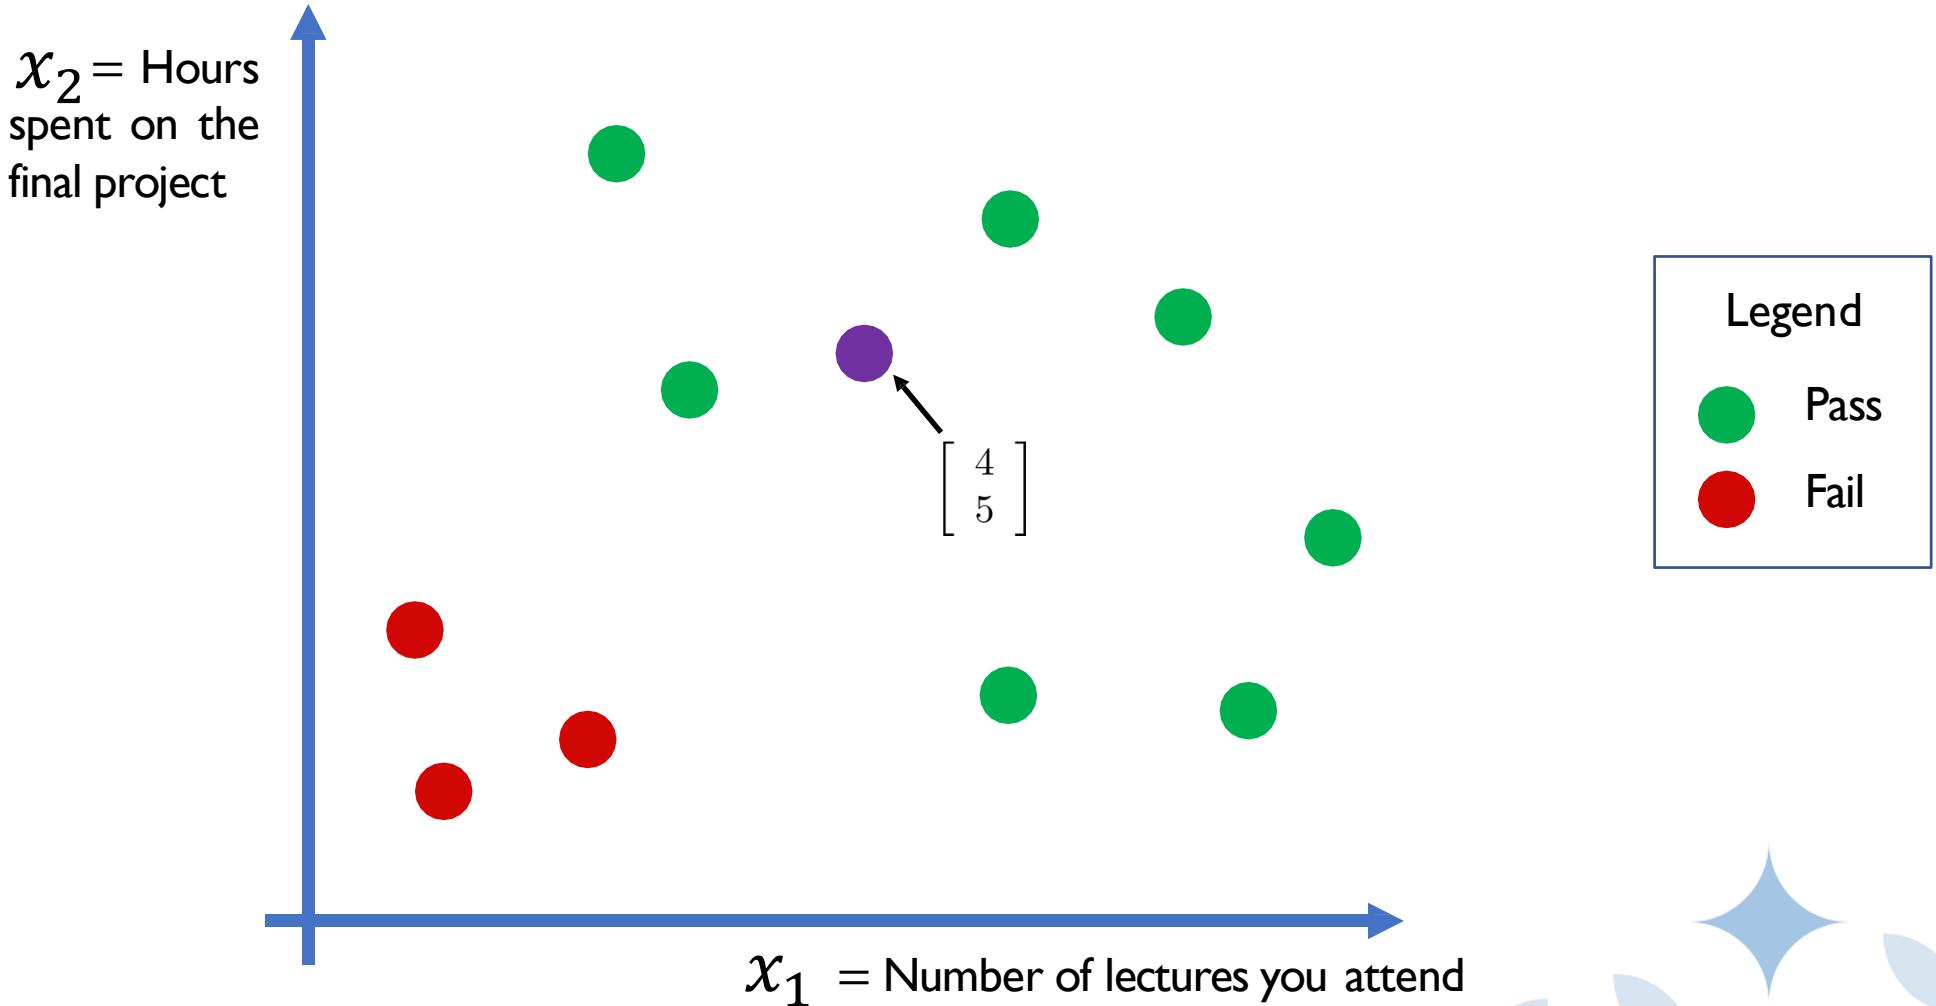
- $x_1$  = Number of lectures you attend
- $x_2$  = Hours spent on the final project



# Example Problem: Will I pass the class?



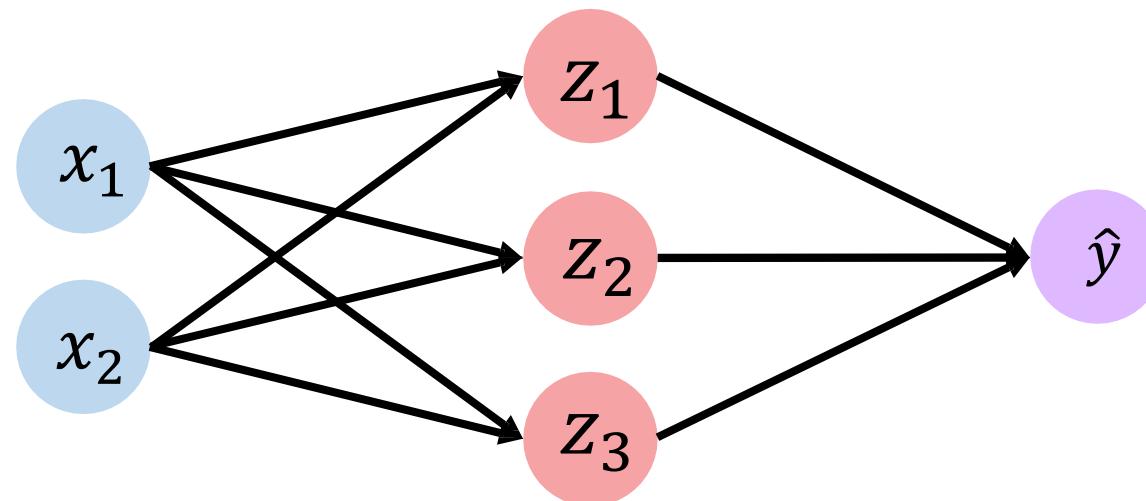
# Example Problem: Will I pass the class?



# Initialization

Initialize the network to realize the first feed-forward pass

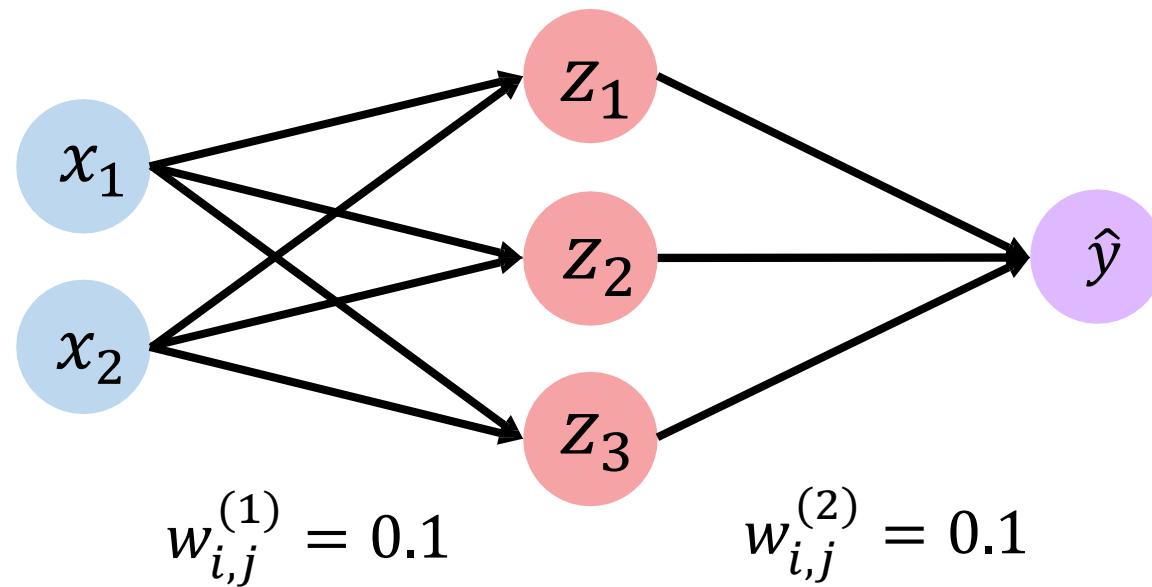
- Step 1: Initialize  $W = (W^{(1)}, W^{(2)})$



# Initialization

Initialize the network for the first feed-forward pass

- Step 1: Initialize  $W = (W^{(1)}, W^{(2)})$

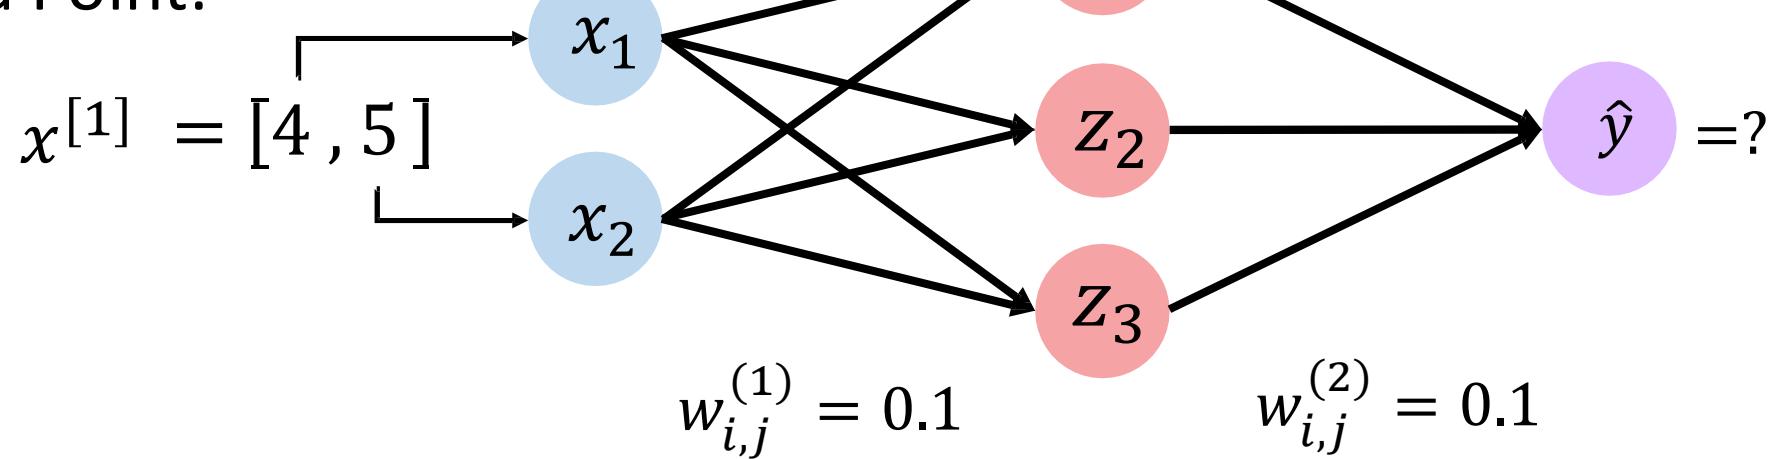


- For example, set all weights to 0.1



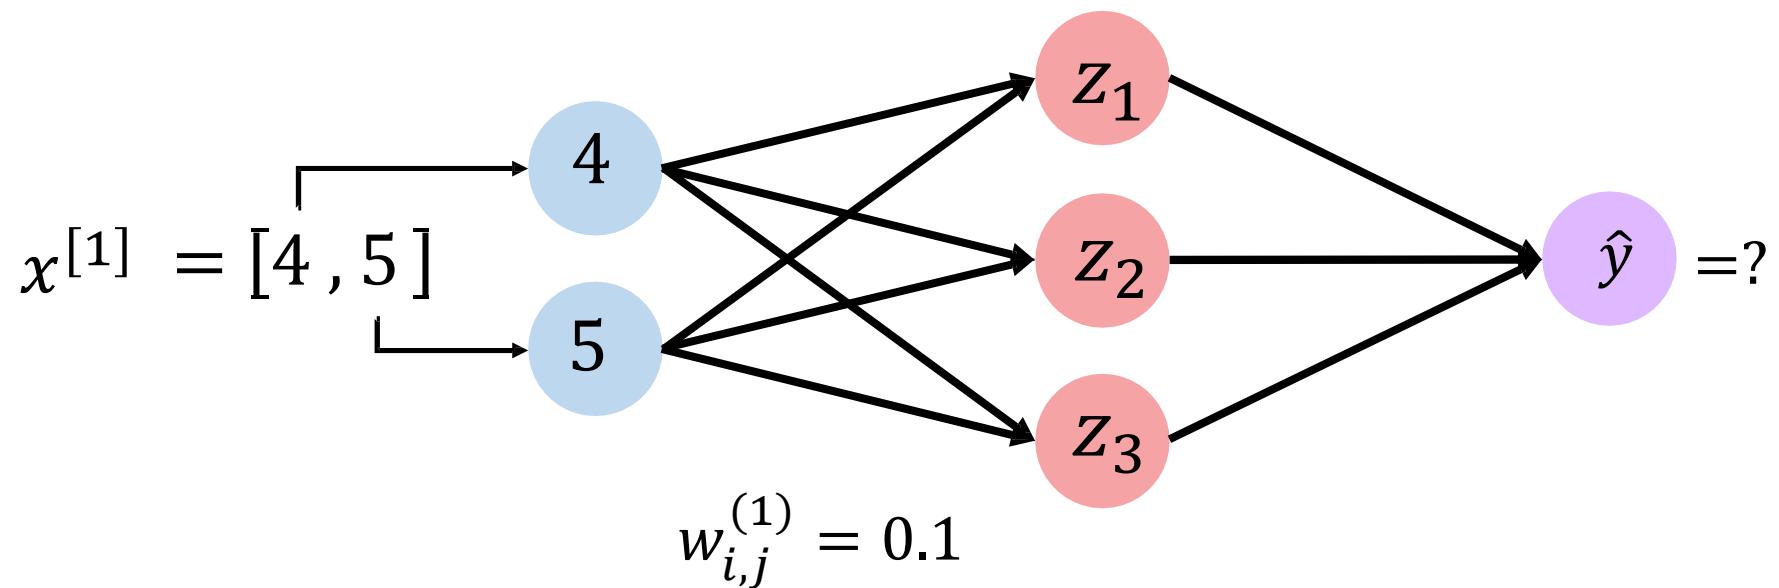
# Feed-forward

Data Point:

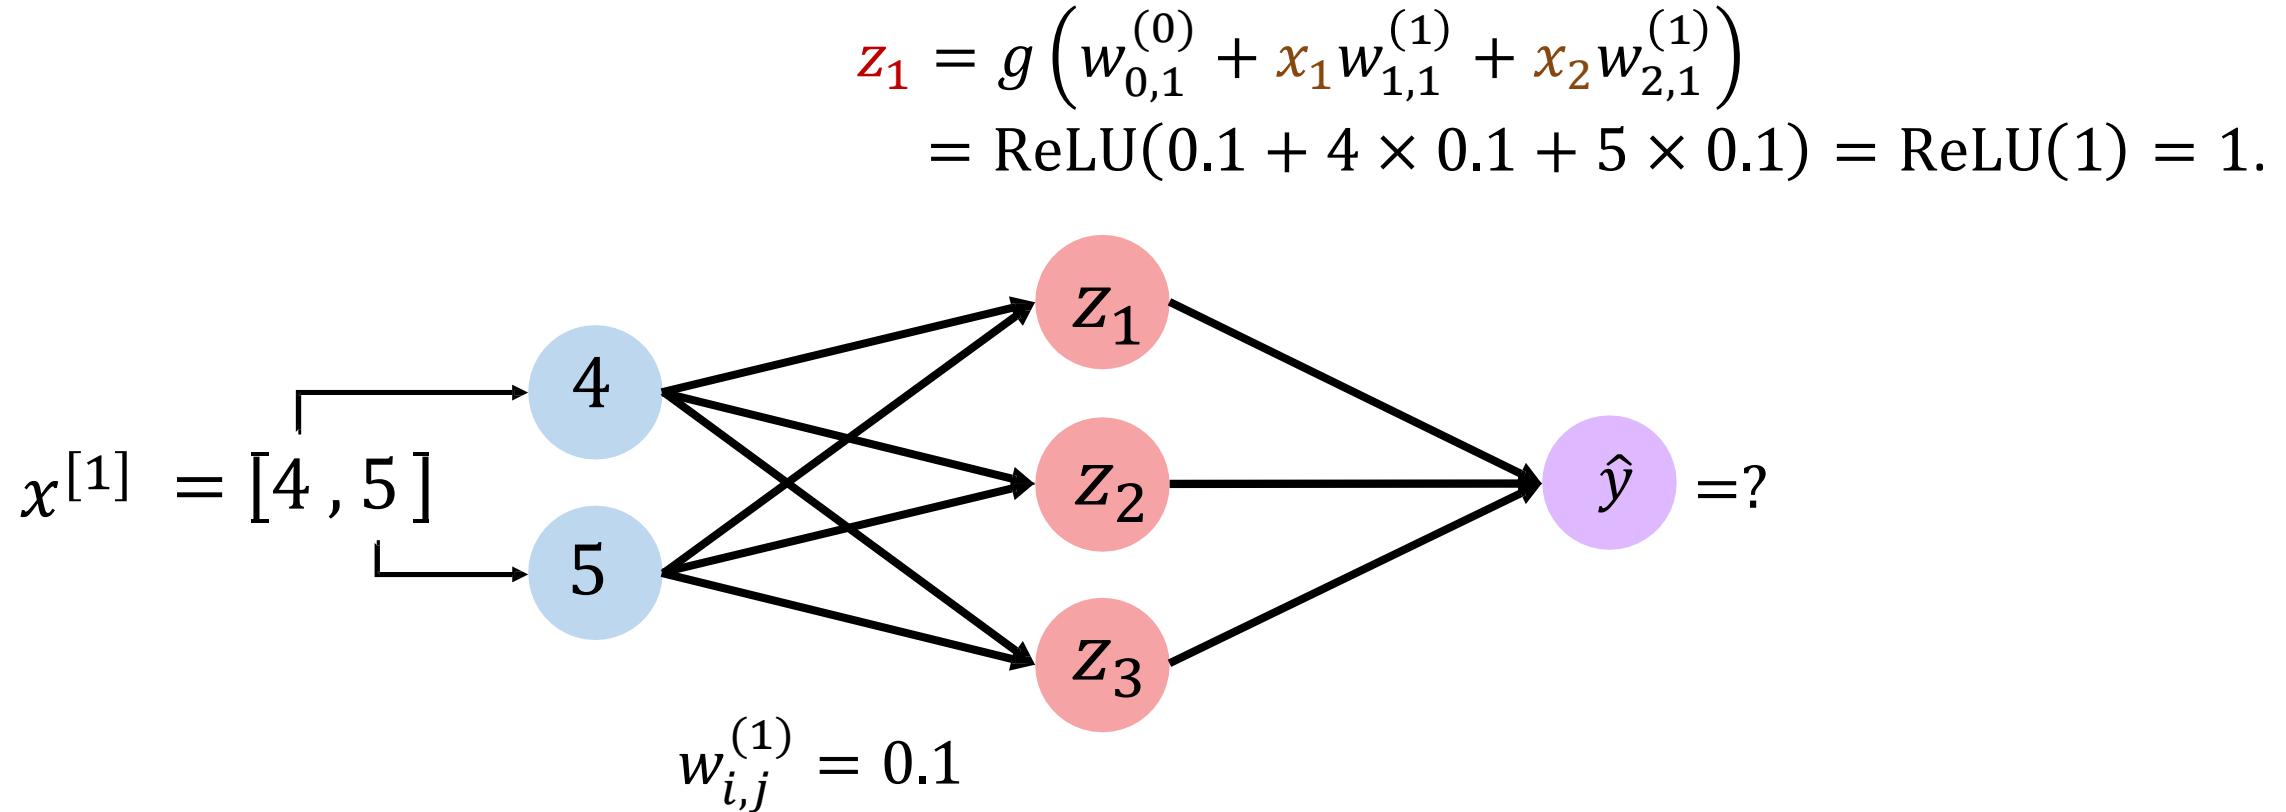


# Feed-forward

$$z_1 = g \left( w_{0,1}^{(0)} + x_1 w_{1,1}^{(1)} + x_2 w_{2,1}^{(1)} \right)$$

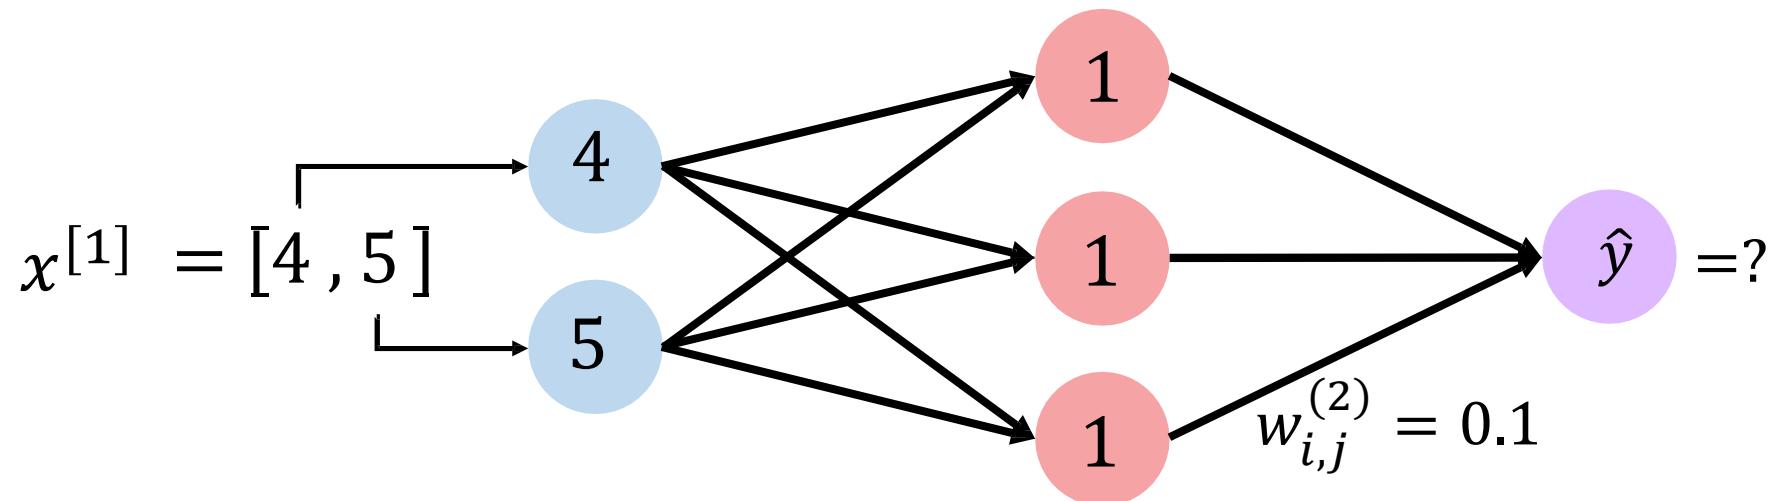


# Feed-forward

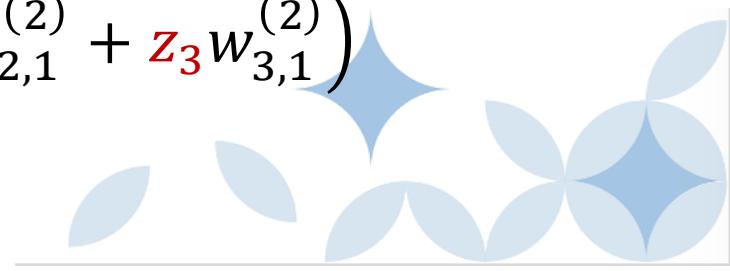


# Feed-forward

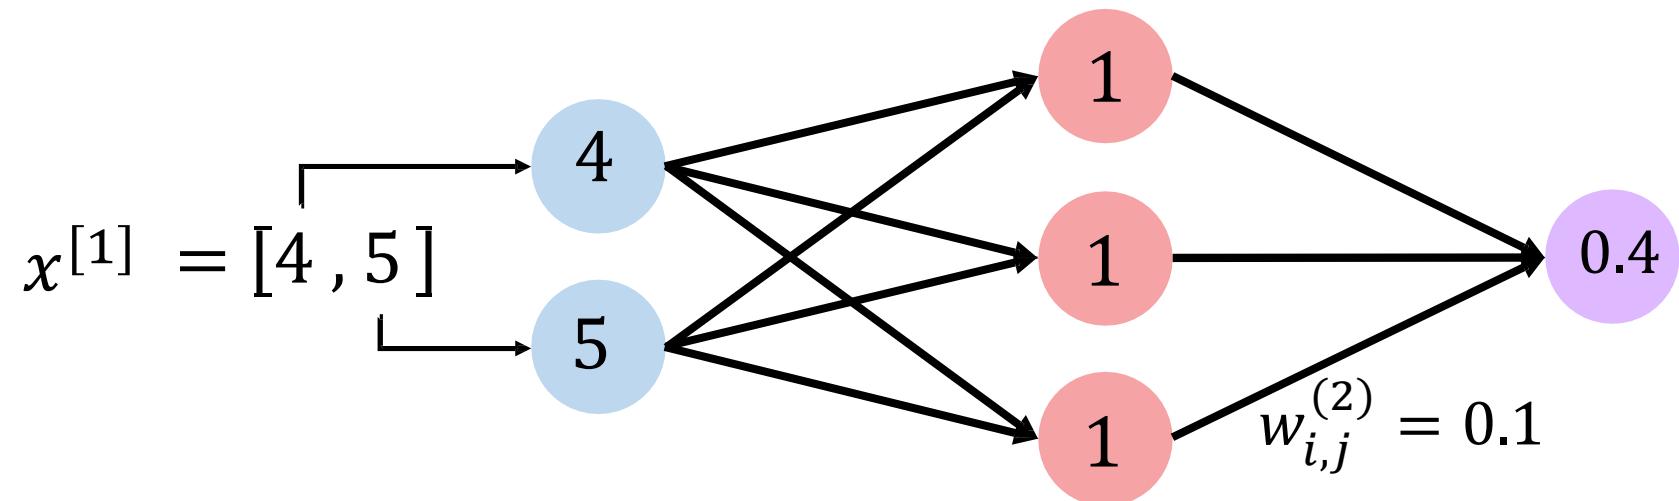
$$\begin{aligned} z_1 &= g\left(w_{0,1}^{(0)} + x_1 w_{1,1}^{(1)} + x_2 w_{2,1}^{(1)}\right) \\ &= \text{ReLU}(0.1 + 4 \times 0.1 + 5 \times 0.1) = \text{ReLU}(1) = 1. \end{aligned}$$



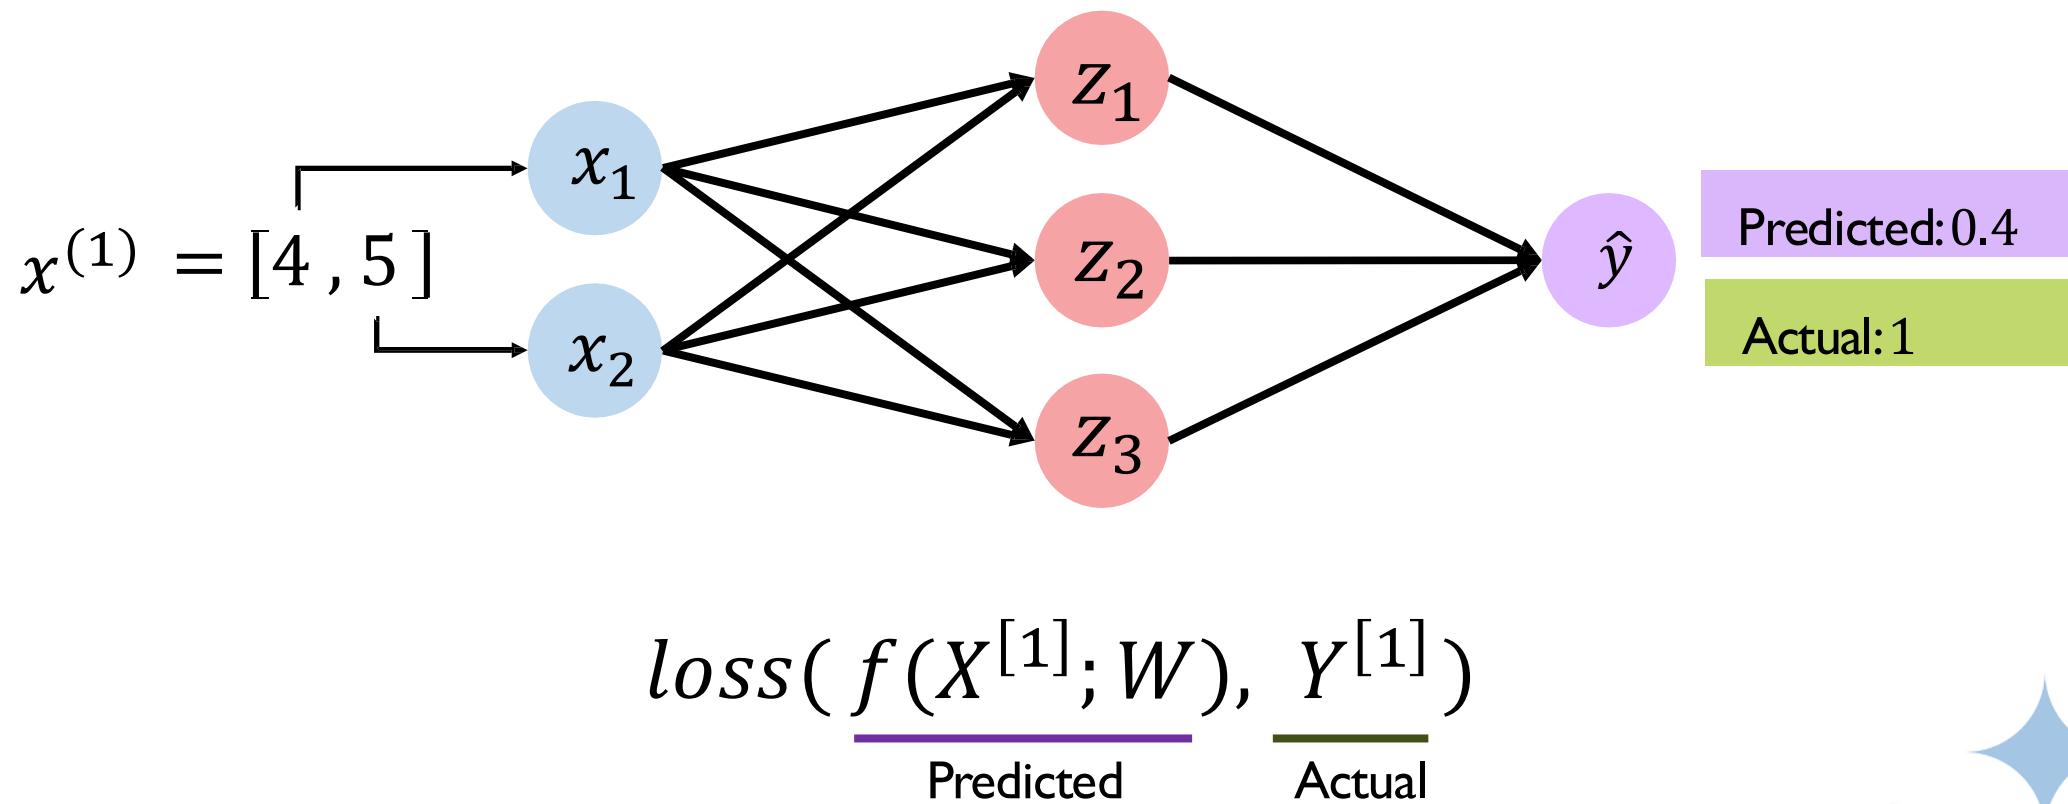
$$\begin{aligned} \hat{y} &= g\left(w_{0,1}^{(2)} + z_1 w_{1,1}^{(2)} + z_2 w_{2,1}^{(2)} + z_3 w_{3,1}^{(2)}\right) \\ &= \text{ReLU}(0.4) = 0.4 \end{aligned}$$



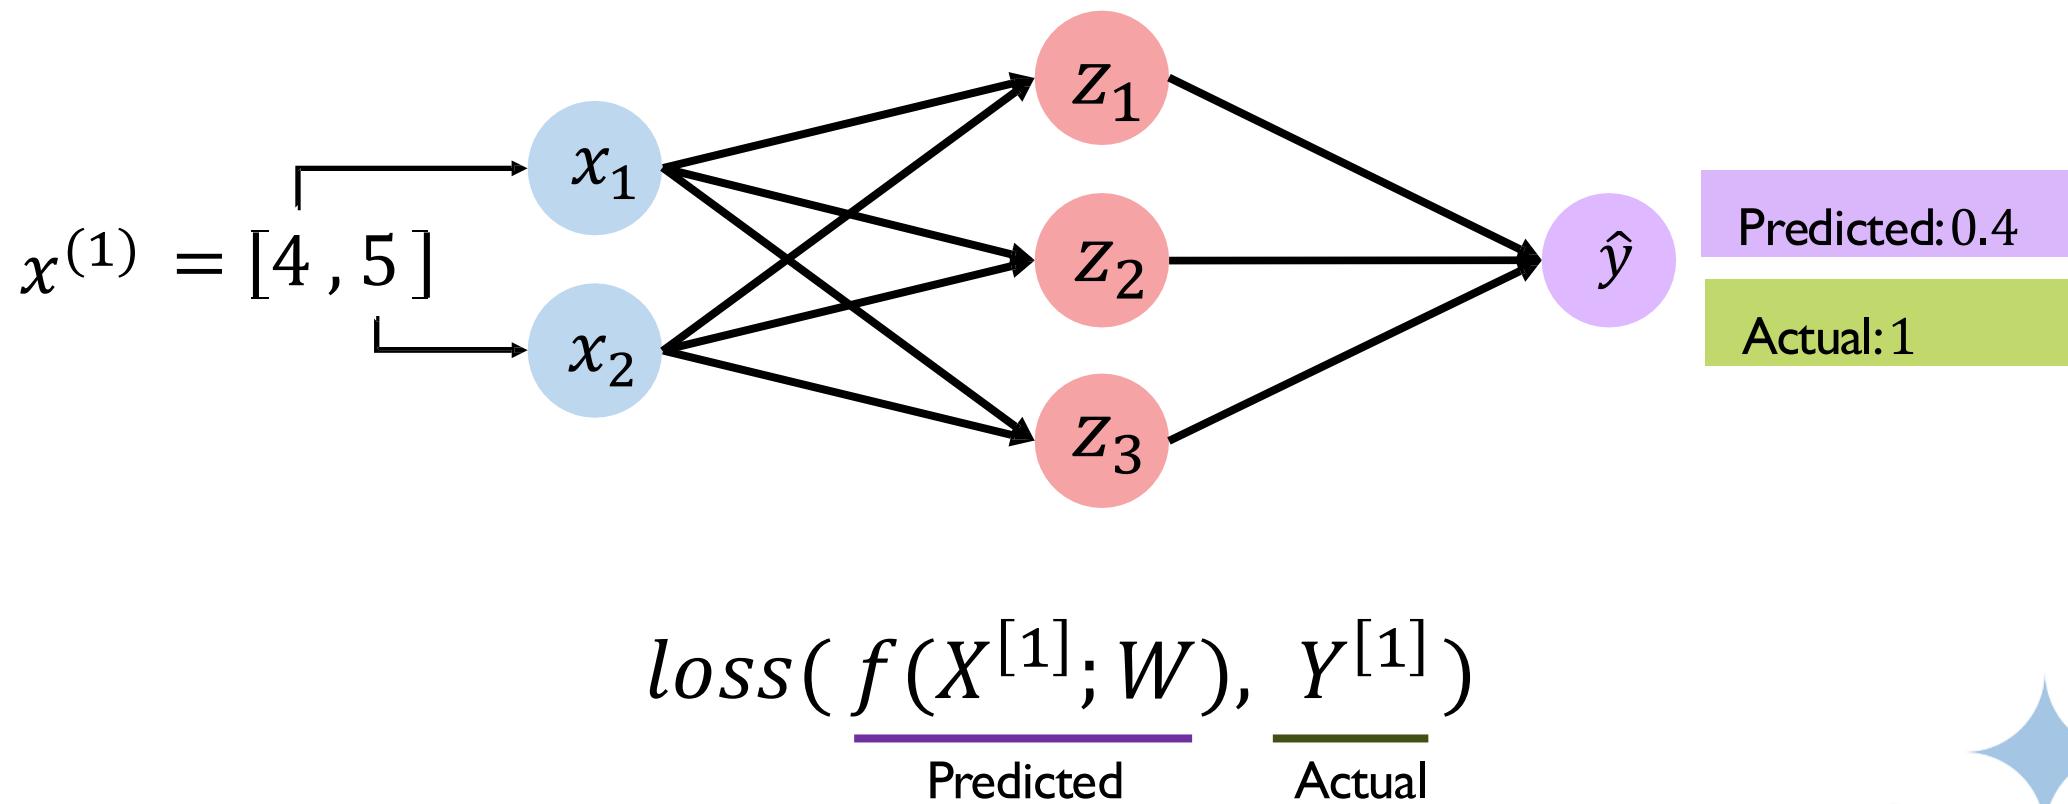
# Feed-forward



# Model Assessment - Quantifying Loss

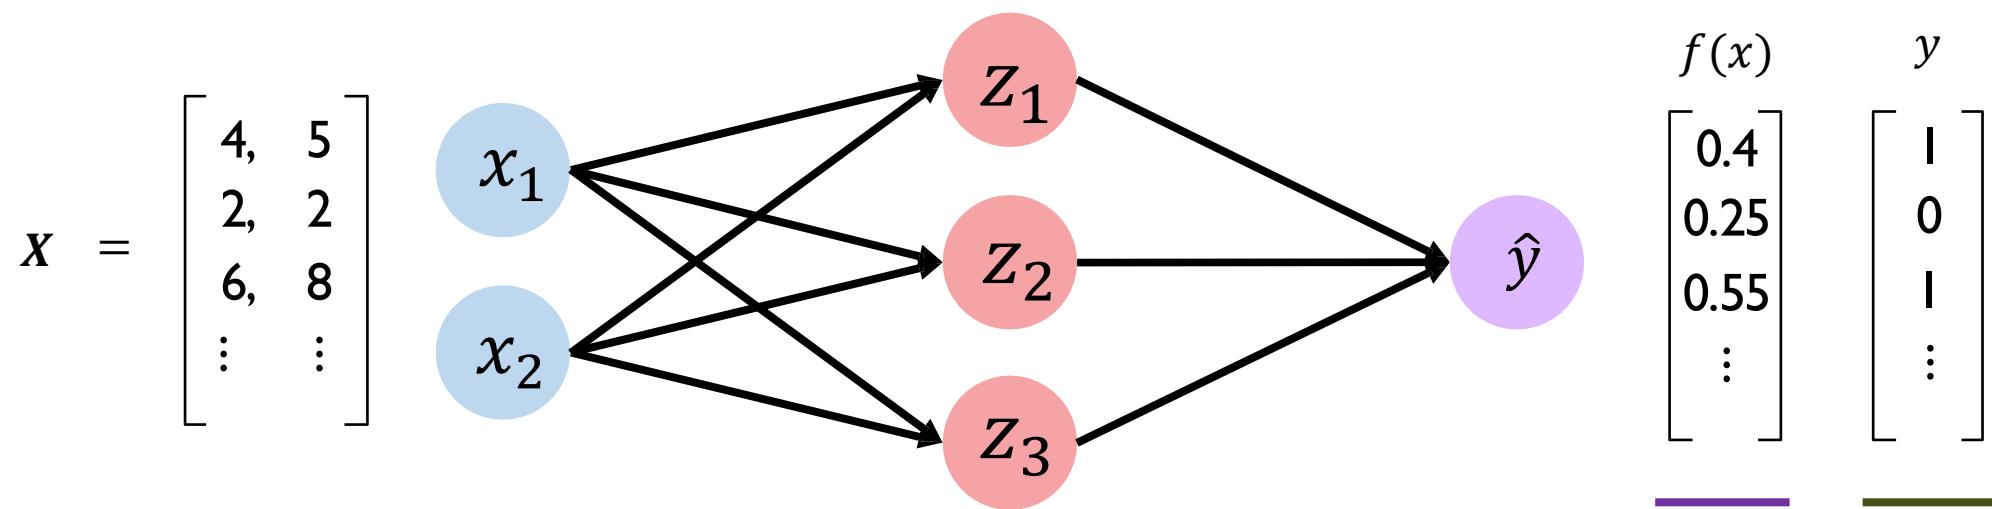


# Model Assessment - Quantifying Loss



# Model Assessment – Empirical Risk

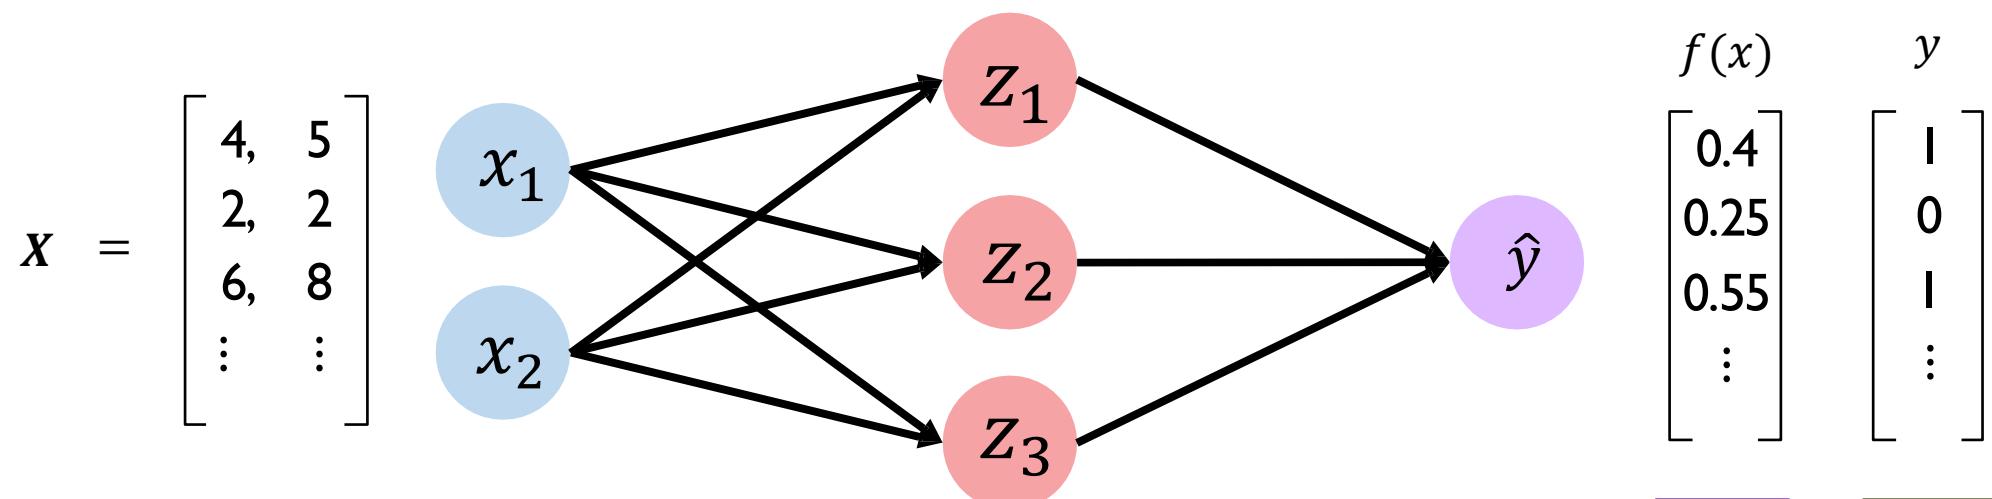
The empirical Risk measures the total loss over the training dataset



$$\min J(W) = \frac{1}{n} \sum_{i=1}^n loss(\underbrace{f(X^{[i]}; W)}_{\text{Predicted}}, \underbrace{Y^{[i]}}_{\text{Actual}})$$

# Model Assessment – Empirical Risk

The empirical Risk measures the total loss over the training dataset



- Also known as:
    - Objective function
    - Cost function

$$\min_{\text{ion}} \overbrace{J(W)}^{\leftarrow} = \frac{1}{n} \sum_{i=1}^n loss(\underbrace{f(X^{[i]}; W)}_{\text{Predicted}}, \underbrace{Y^{[i]}}_{\text{Actual}})$$

# Quantifying Loss – Binary Classification

- 0/1 Loss:  $loss(Y, f(X)) = 1_{Y \neq f(X)}$ 
  - $loss = 0$  when correctly classified, 1 when misclassified
  - Intuitive but hard to train

$$\min J(W) = \frac{1}{n} \sum_{i=1}^n loss(f(X^{[i]}; W), \underbrace{Y^{[i]}}_{\text{Actual}})$$

Predicted



# Quantifying Loss – Binary Classification

- 0/1 Loss:  $loss(Y, f(X)) = 1_{Y \neq f(X)}$ 
  - $loss = 0$  when correctly classified, 1 when misclassified
  - Intuitive but hard to train
- Binary Cross Entropy Loss:  $loss(Y, f(X)) = -Y \log f(X) - (1 - Y) \log(1 - f(X))$ 
  - measure of the difference between two probability distributions
  - Can be used when output a probability from 0 to 1

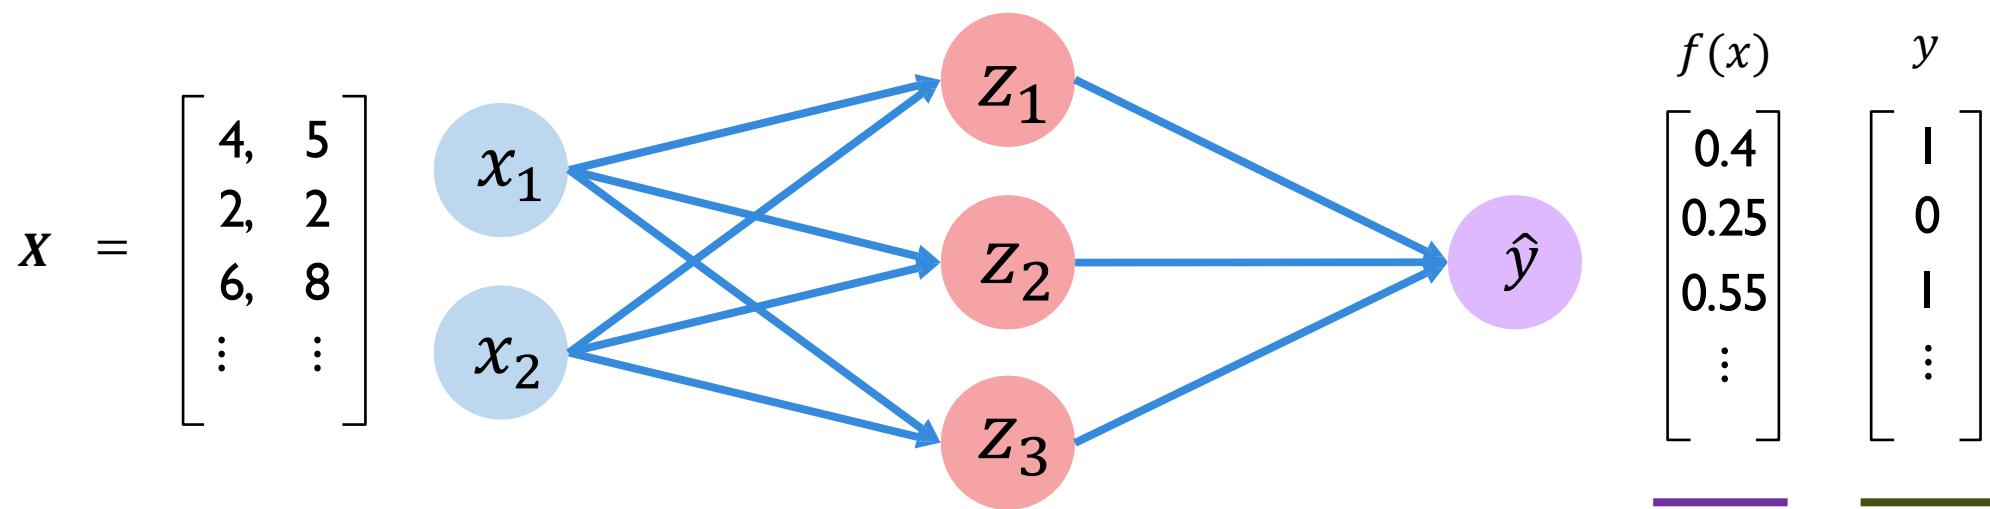
$$\min J(W) = \frac{1}{n} \sum_{i=1}^n loss(f(X^{[i]}; W), \underbrace{Y^{[i]}}_{\text{Actual}})$$

Predicted



# Classification: Binary Cross Entropy Loss

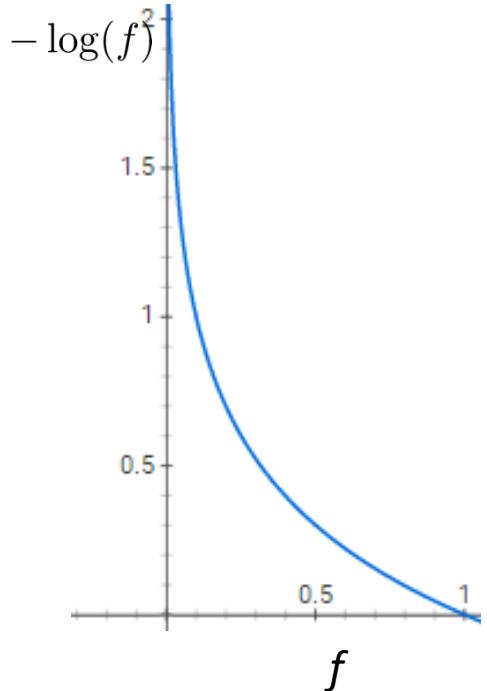
Using Binary Cross Entropy Loss for binary classification



$$\min J(W) = \frac{1}{n} \sum_{i=1}^n \underbrace{-Y^{[i]} \log f(X^{[i]}; W)}_{\text{Actual}} - \underbrace{(1 - Y^{[i]}) \log(1 - f(X^{[i]}; W))}_{\text{Predicted}}$$

# Classification: Binary Cross Entropy Loss

True label  
 $Y^{[i]} = 1$



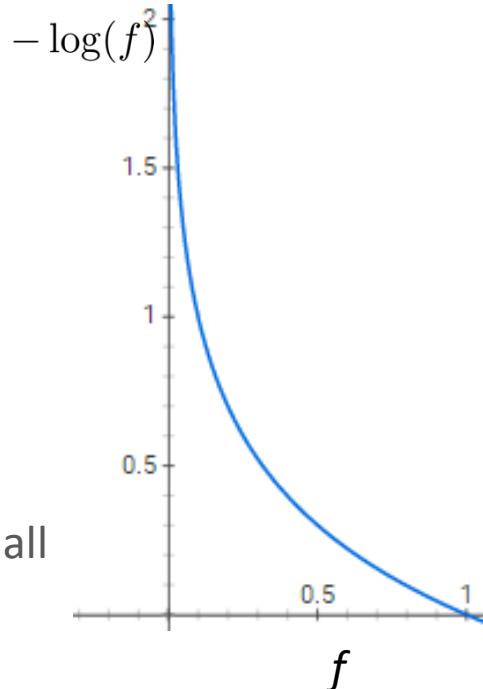
$$\min J(W) = \frac{1}{n} \sum_{i=1}^n \underbrace{-Y^{[i]} \log f(X^{[i]}; W)}_{\text{Actual}} - \underbrace{(1 - Y^{[i]}) \log(1 - f(X^{[i]}; W))}_{\text{Predicted}}$$

# Classification: Binary Cross Entropy Loss

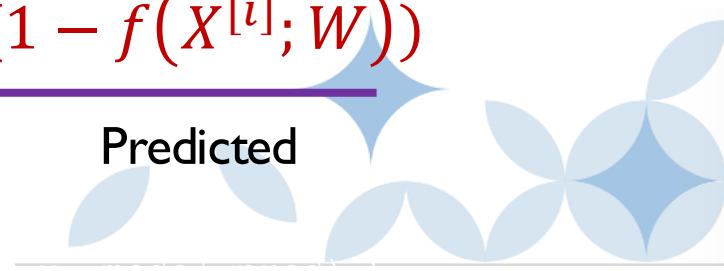
True label  
 $Y^{[i]} = 1$



Hope  $-\log(f)$  to be small  
such that  $f$  is close to 1



$$\min J(W) = \frac{1}{n} \sum_{i=1}^n \underbrace{-Y^{[i]} \log f(X^{[i]}; W)}_{\text{Actual}} - \underbrace{(1 - Y^{[i]}) \log(1 - f(X^{[i]}; W))}_{\text{Predicted}}$$

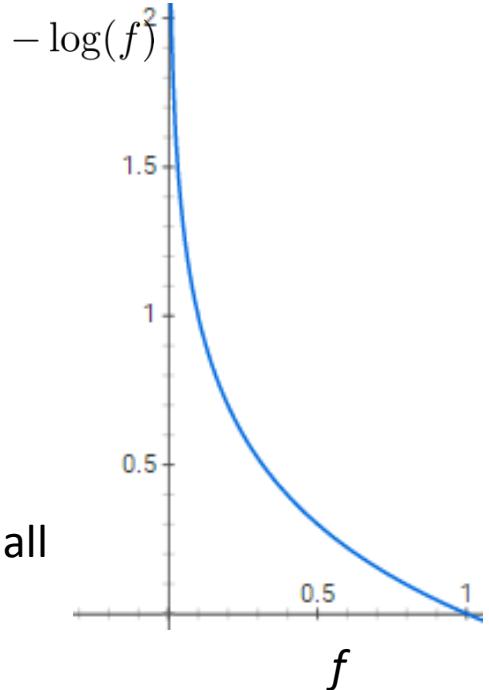


# Classification: Binary Cross Entropy Loss

True label  
 $Y^{[i]} = 1$



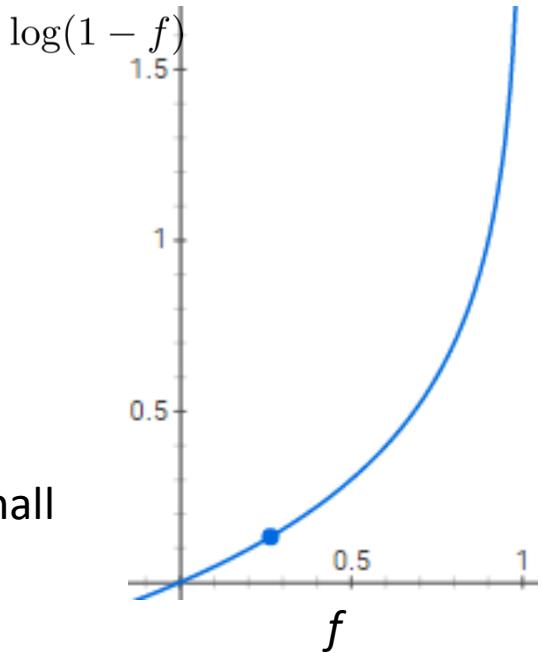
Hope  $-\log(f)$  to be small  
such that  $f$  is close to 1



True label  
 $Y^{[i]} = 0$



Hope  $-\log(1 - f)$  to be small  
such that  $f$  to be close to 0



$$\min J(W) = \frac{1}{n} \sum_{i=1}^n \underbrace{-Y^{[i]} \log f(X^{[i]}; W)}_{\text{Actual}} + \underbrace{(1 - Y^{[i]}) \log(1 - f(X^{[i]}; W))}_{\text{Predicted}}$$

Predicted

# Classification: Binary Cross Entropy Loss

- Binary Cross Entropy Loss

$$\text{loss} = -Y \log f(X; W) - (1 - Y) \log(1 - f(X; W))$$

- Vectorize True Label as  $(Y_1, Y_0)$ :  $Y_1 = Y$  and  $Y_0 = 1 - Y$

- Prediction:  $\hat{Y}_1 = f(X; W)$  versus  $\hat{Y}_0 = 1 - f(X; W)$

$$\text{loss} = -Y_1 \log \hat{Y}_1 - Y_0 \log \hat{Y}_0$$

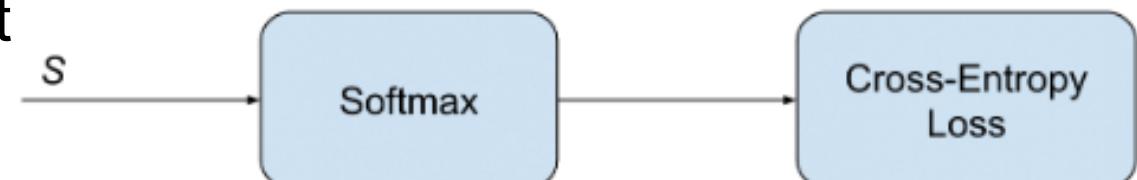
- Note that  $Y_1 + Y_0 = 1$  and  $\hat{Y}_1 + \hat{Y}_0 = 1$

- Extend the definition Multi-Class?



# Classification: Multi-class Cross Entropy Loss

- Binary Cross-entropy  $loss = -Y_1 \log \hat{Y}_1 - Y_0 \log \hat{Y}_0$
- $K$  Classes:  $loss = -\sum_{k=1,\dots,K} Y_k \log(\hat{Y}_k)$
- Vectorize  $Y = (Y_1, Y_2, Y_3, \dots)^\top$ :  $loss = -Y^T \log(\hat{Y})$
- The true label vector  $\sum_k y_k = 1$ . How to make sure  $\sum_k \hat{y}_k = 1$ ?
  - Apply a **Softmax** operator to the output
  - Use it as the activation function of the **last** layer
  - When no hidden layer: logistic regression



$$\hat{y}_k = \frac{\exp(s_k)}{\sum \exp(s_i)}$$

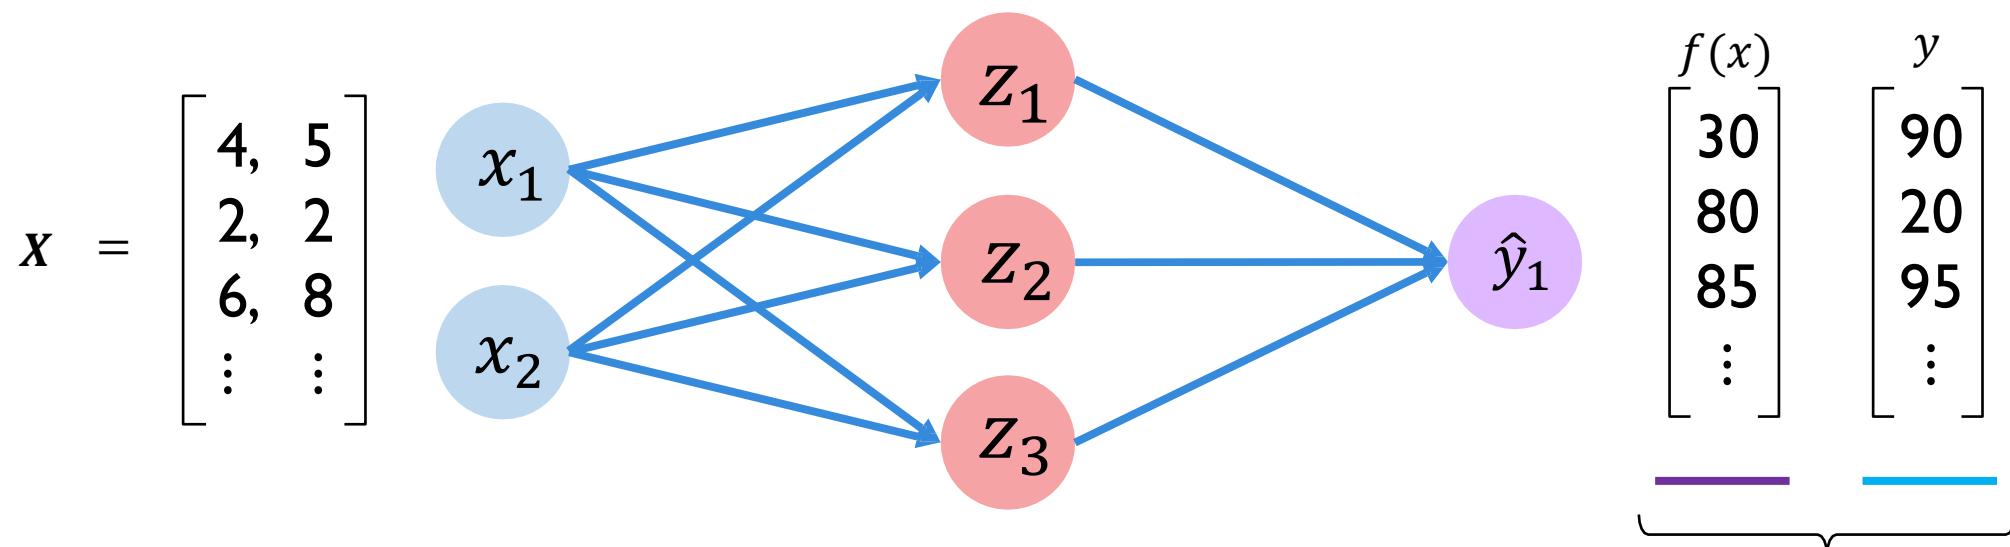


# Regression: Mean Squared Error Loss

**Used with regression models that output continuous real numbers**

# Regression: Mean Squared Error Loss

Used with regression models that output continuous real numbers



$$\min J(W) = \frac{1}{n} \sum_{i=1}^n \left( \underbrace{f(X^{[i]}; W)}_{\text{Predicted}} - \underbrace{Y^{[i]}}_{\text{Actual}} \right)^2$$

# Loss Optimization

We want to find the network weights that achieve the lowest loss

$$W^* = \underset{W}{\operatorname{argmin}} J(W) = \frac{1}{n} \sum_{i=1}^n \text{loss}(f(X^{[i]}; W), Y^{[i]})$$

↑  
 $W = \{W^{(1)}, W^{(2)}, \dots\}$

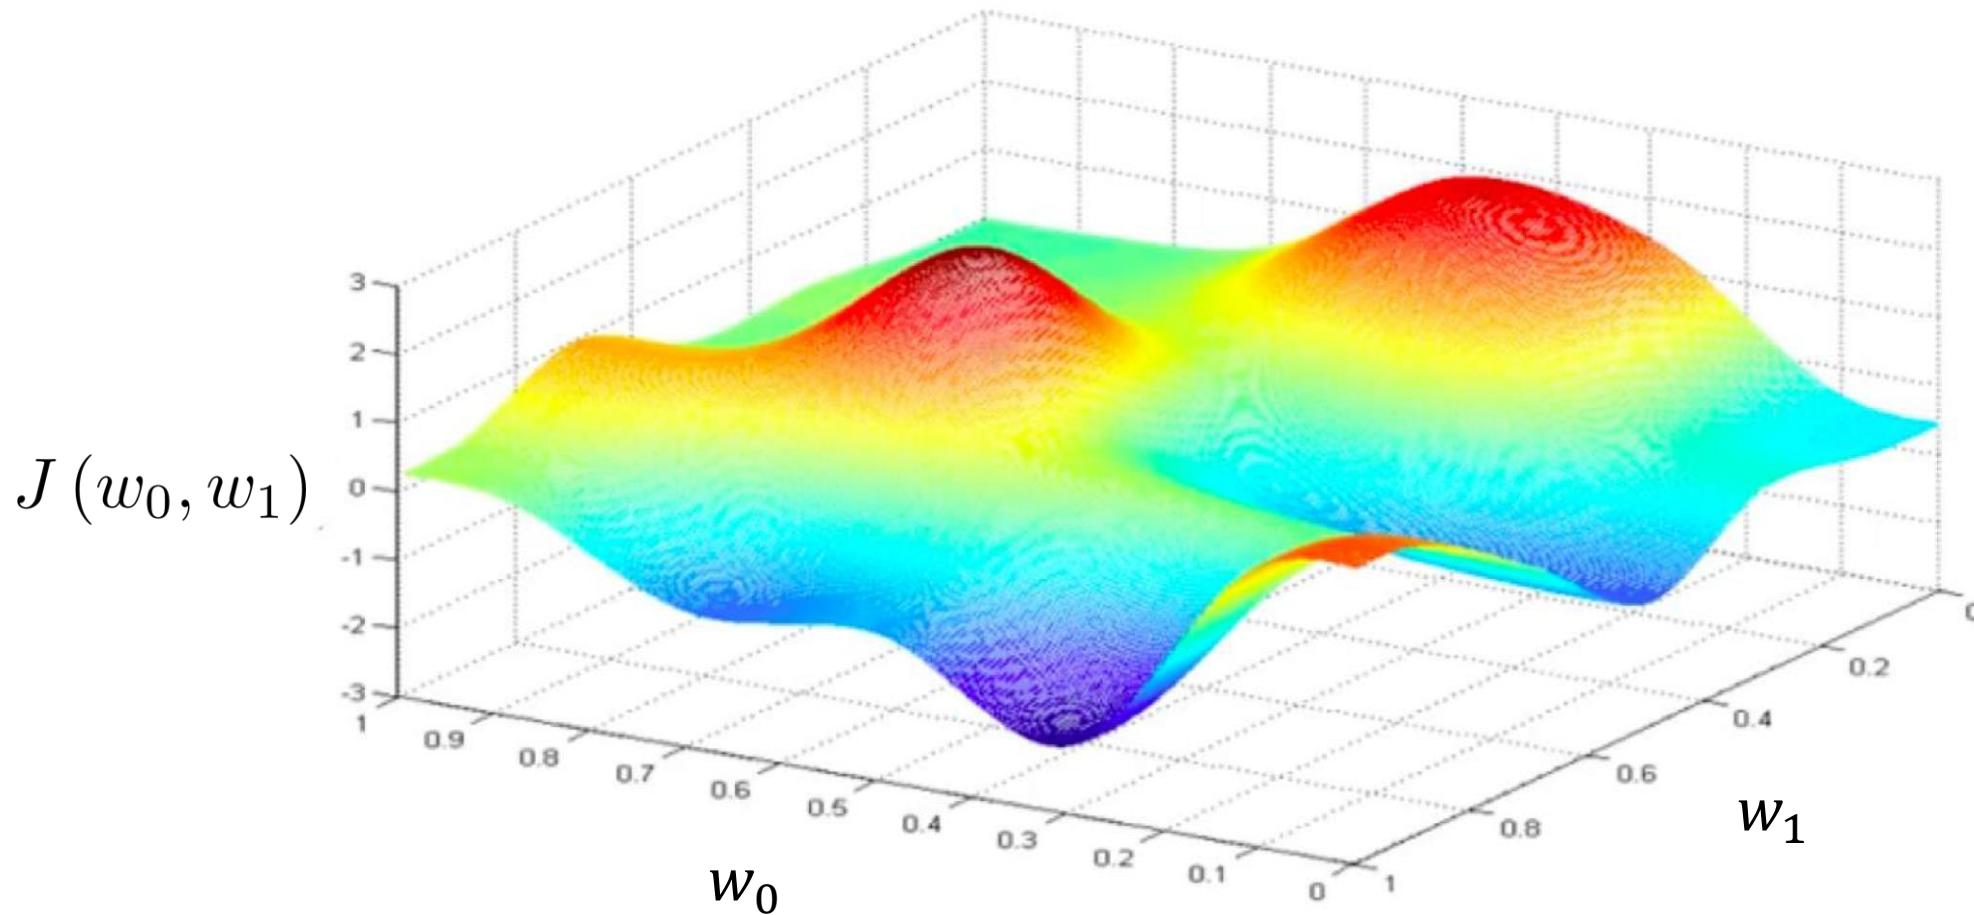
Remember:

*Our loss is a function of the network weights!*



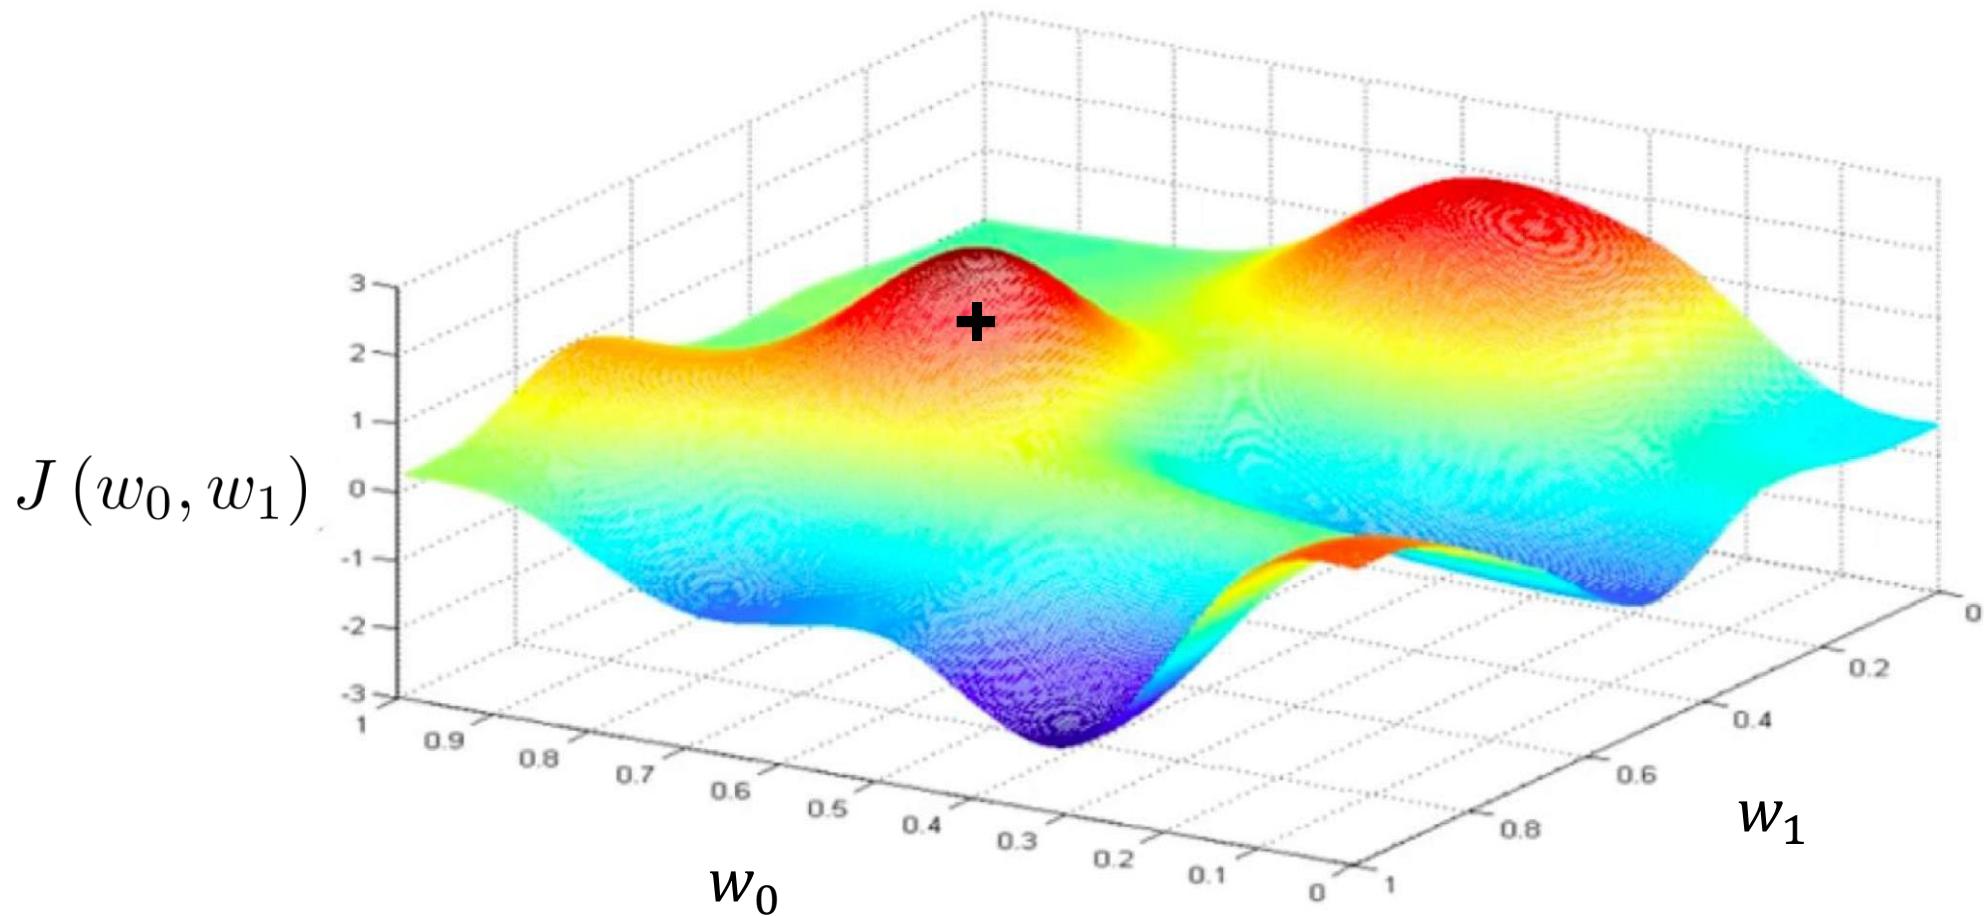
# Loss Optimization - Gradient Descent

$$W^* = \underset{W}{\operatorname{argmin}} J(W)$$



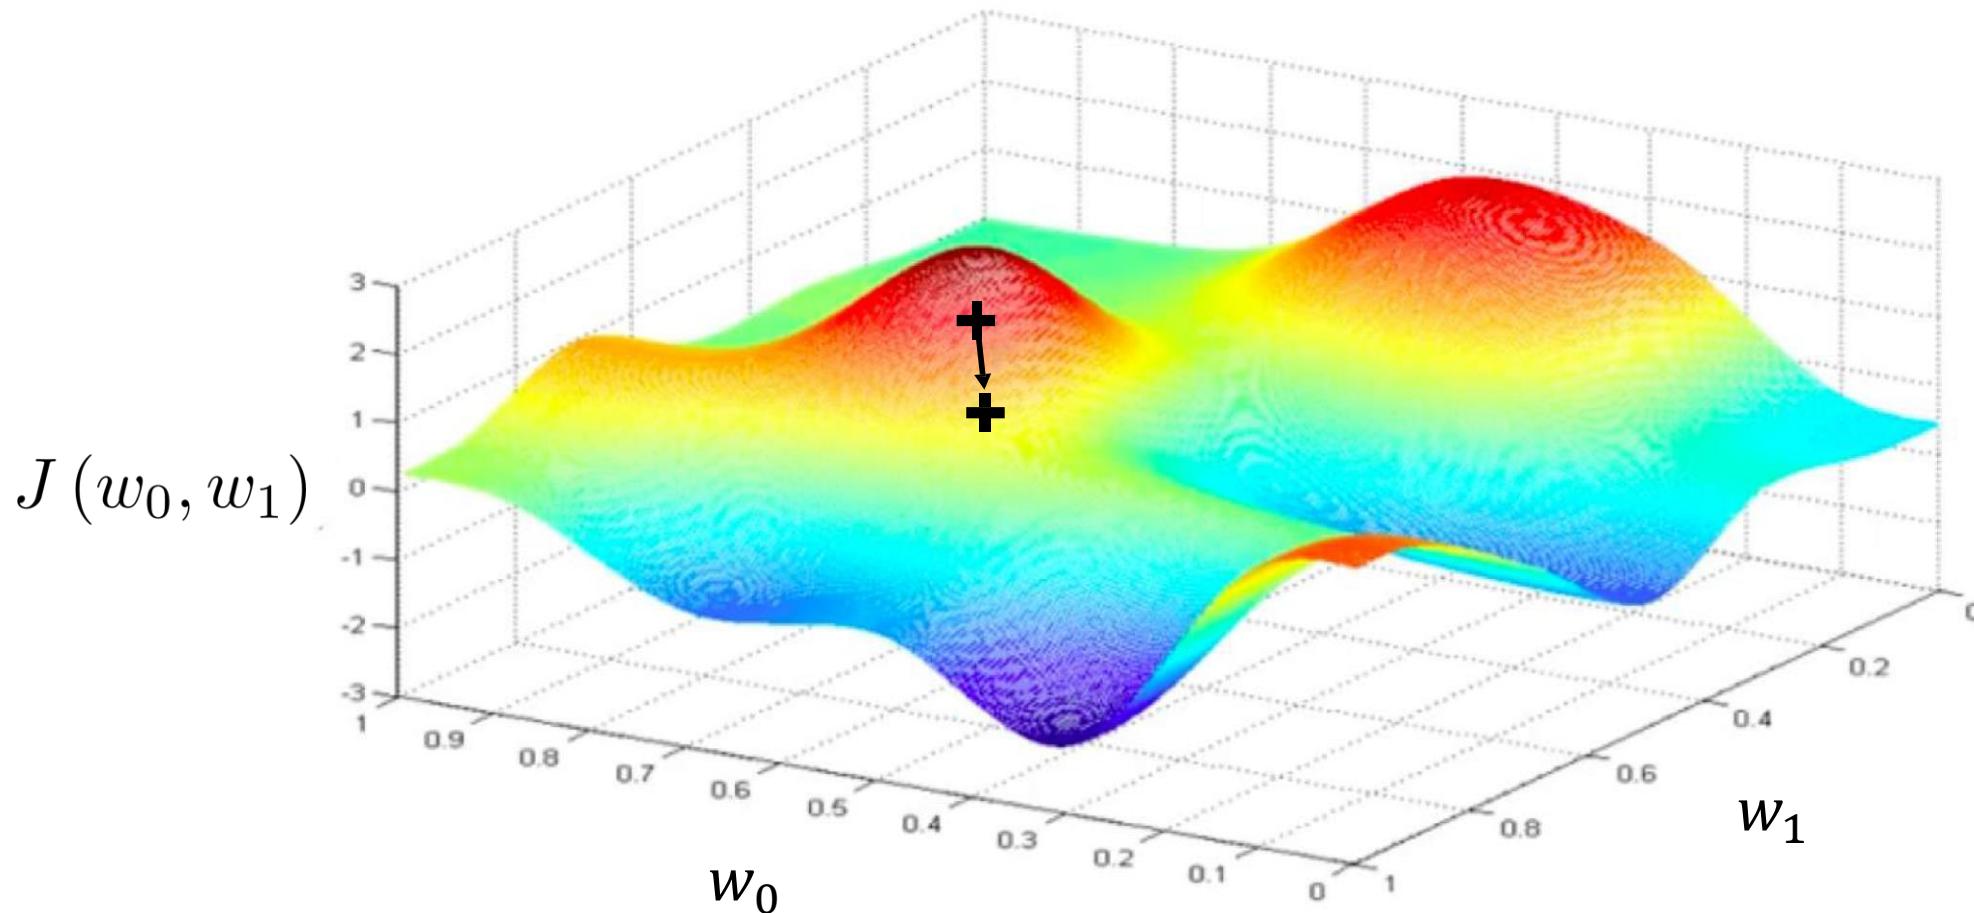
# Loss Optimization - Gradient Descent

Randomly pick an initial  $(w_0, w_1)$



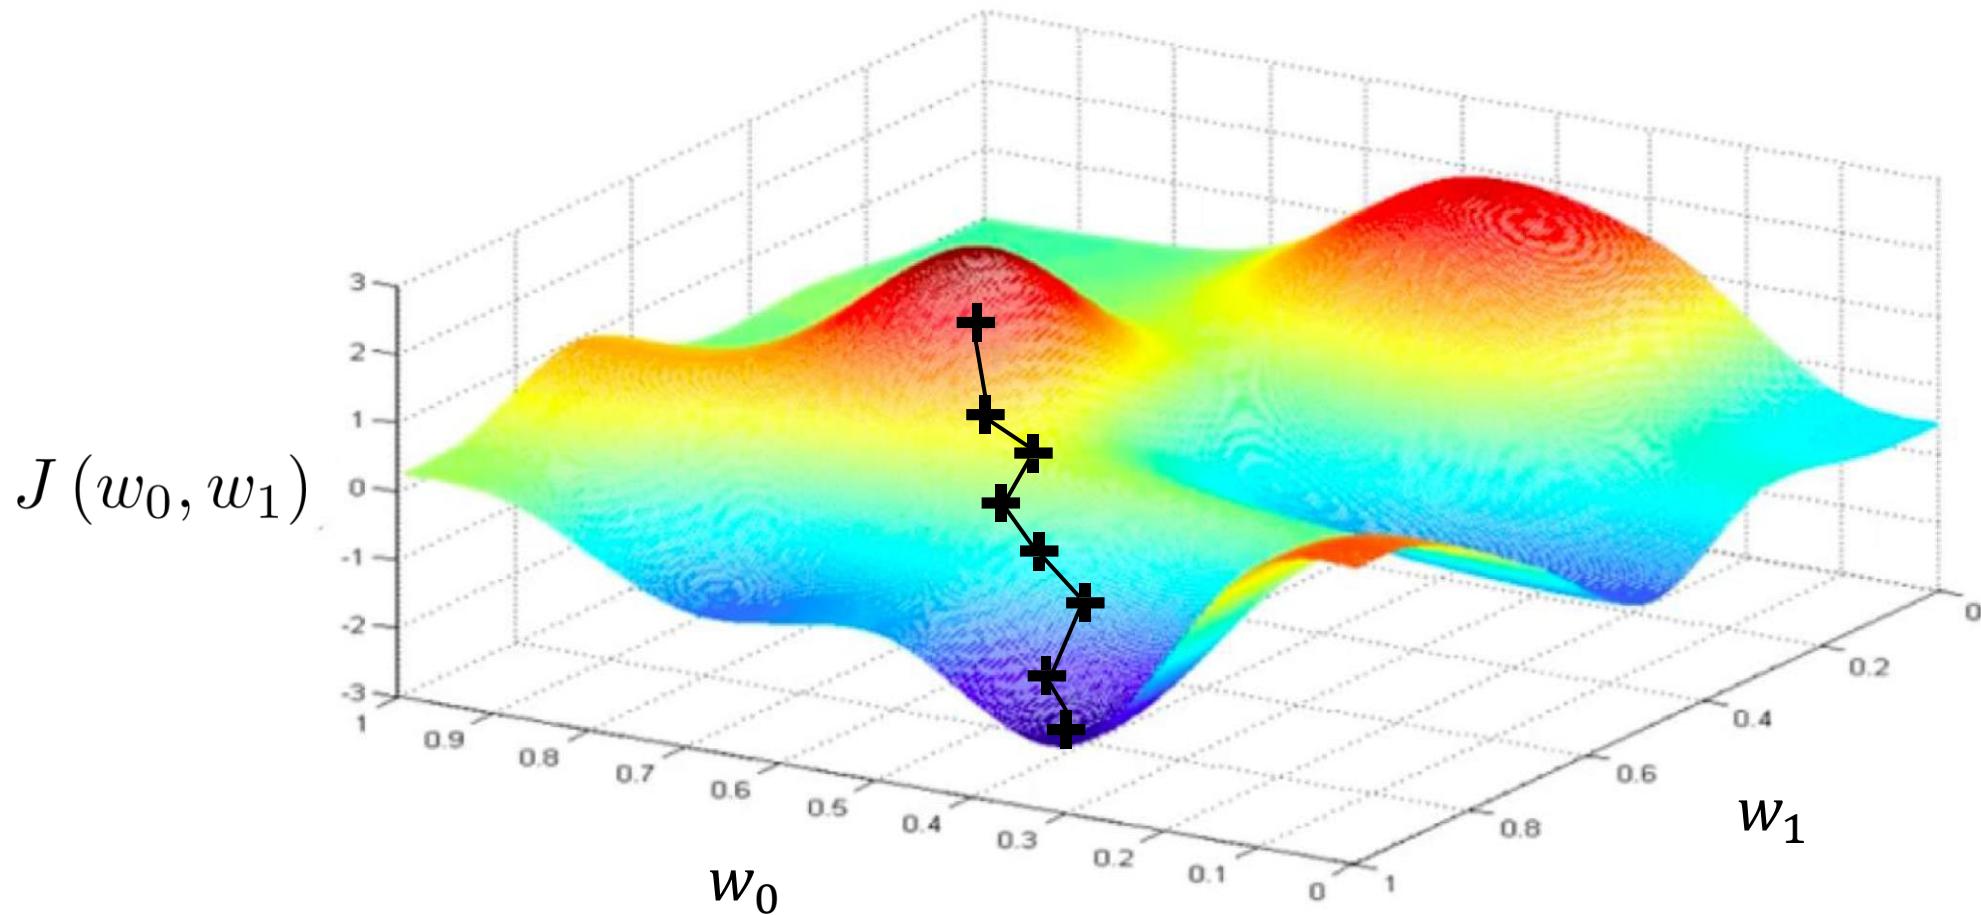
# Loss Optimization - Gradient Descent

Compute gradient and move towards negative gradient direction



# Loss Optimization - Gradient Descent

Repeat until convergence



# Loss Optimization - Gradient Descent

1. Initialize weights
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$
4. Update weights,  $W^{t+1} \leftarrow W^t - \eta \frac{\partial J(W)}{\partial W^\top}$
5. Return weights



# Loss Optimization - Gradient Descent

1. Initialize weights Question 1: How to set initial weights?
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$
4. Update weights,  $W^{t+1} \leftarrow W^t - \eta \frac{\partial J(W)}{\partial W^\top}$
5. Return weights



# Loss Optimization - Gradient Descent

1. Initialize weights      Question 1: How to set initial weights?
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$       Question 2: How to compute gradient?
4. Update weights,  $W^{t+1} \leftarrow W^t - \eta \frac{\partial J(W)}{\partial W^\top}$
5. Return weights



# Loss Optimization - Gradient Descent

1. Initialize weights

Question 1: How to set initial weights?

2. Loop until convergence:

3. Compute gradient,

$$\frac{\partial J(W)}{\partial W}$$

Question 2: How to compute gradient?

4. Update weights,  $W^{t+1} \leftarrow W^t - \eta \frac{\partial J(W)}{\partial W^\top}$

Question 3: How to choose the stepsize?

5. Return weights

$\eta$ : Stepsize (Learning Rates)



# Loss Optimization - Gradient Descent

1. Initialize weights Question 1: How to set initial weights?
  2. Loop until convergence: Question 4: When to stop?
  3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$  Question 2: How to compute gradient?
  4. Update weights,  $W^{t+1} \leftarrow W^t - \eta \frac{\partial J(W)}{\partial W^T}$  Question 3: How to choose  $\eta$ ?  
5. Return weights  $\eta$ : Stepsize (Learning Rates)

$\eta$ : Stepsize (Learning Rates)

Thanks!



Questions?